Crawler Access and Content Effects in Answer Engine Optimization: A Website Audit and Controlled Experiments

Ayush Agarwal
Independent researcher ayushag1002@gmail.com

November 2026

Abstract

Answer engines such as ChatGPT and Perplexity name organisations and cite sources. We study whether crawlers can reach a site's own pages and how page content affects selection. Study 1 audits 569 organisational websites in 22 sectors and 703 popular domains. Of the organisational sites, 35.0% showed at least one potential obstacle on a first-party crawling route (19.0% without research-client refusals at the network edge), mostly not expressed in agent-specific robots.txt rules, and sites restricted AI training crawlers far more often than answer agents; these are measured obstacles, not demonstrated exclusion. Study 2 is a pre-specified experiment with fictional brands: 14 page interventions in 12 hosted API models, replicated on 11 more. With pages supplied, a third party's recommendation raised the chance that a brand was named by 37.4 percentage points, more than any edit to the brand's own page; citation and statistics packages added 21.1 and 13.7 points, keyword stuffing also helped, and an instruction embedded in the page worked on most models but not on Anthropic's. In a follow-up, recommending rather than neutral wording on the brand's page added 10.1 points, and crediting an outside reviewer 3.1 more. When four models searched a synthetic corpus themselves, the brand was named in 22% of baseline answers instead of 53%, evidence and instruction effects were no longer detectable, and text matching the question made pages easier to find. Outcomes record whether answers name the brand, not whether they recommend it. We release the toolkit, benchmark and data.

1 Introduction

Answer engines such as ChatGPT, Claude, Gemini and Perplexity return written answers that may name organisations and cite sources. Users click through less than they do from a list of links: in a 2025 panel study, they clicked a conventional result on 8% of Google visits that showed an AI summary and on 15% of visits that did not [4]. Being named, recommended and cited has become a goal in its own right, pursued under the names Answer Engine Optimization (AEO) and Generative Engine Optimization (GEO) [1]. An answer may draw on a publisher's own pages, third-party pages, or information already held by a model; these routes are not interchangeable.

Prior work often edits a page already supplied to a model and measures whether it becomes more visible. Results vary across settings: Aggarwal et al. [1] reported gains from statistics, quotations and citations, while C-SEO Bench found weaker effects when many documents compete [36]; other controlled studies emphasise relevance and position over stylistic cues [41, 42]. On real pages, an intervention can also change the information provided, and many add text, so content and length are difficult to separate.

Access conditions are another part of the process. Operators distinguish agents used for model training, search indexing and user-requested fetches, each with its own robots.txt token (Table 1).

Content-delivery networks filter automated clients, and the largest has blocked AI crawlers by default for new customers since July 2025 [6]; most AI crawlers do not execute JavaScript [40]. Prior measurements have examined controls on AI training data [9, 25, 28]. We measure potential obstacles along selected first-party access routes and test content interventions under two controlled task regimes, without treating an observed obstacle as demonstrated exclusion from answers.

We ask three questions, following the stages of the decomposition in Section 3:

RQ1. *Access.* How often do organisations' websites present potential obstacles to the crawlers that answer engines use, and of what kinds?

RQ2. *Selection.* When a model has a brand's pages in its context, which changes to those pages, or to a third-party page, change whether the model names the brand?

RQ3. *Search.* Do those effects persist when the model must search for the pages itself?

Study 1 (Section 4) addresses RQ1 with an audit of 569 purposively selected organisational websites in 22 sectors and a separate, rank-stratified Tranco sample of 703 popular domains. Study 2 (Section 5) addresses RQ2 with a pre-specified experiment that places fictional brands in the context of 12 hosted API models from ten developers and applies 14 page interventions, including a boilerplate placebo and an instruction embedded in the page; a replication on 11 additional API models, three small local models and robustness checks test how far the results carry, and two follow-ups separate a recommendation's wording from its stated source and from the page that carries it. A searchable-corpus task with four models addresses RQ3, and single-account observations of three consumer assistants record whether they visibly searched and what they cited. These designs estimate effects under the tested conditions, not returns in production engines.

The paper makes three contributions: (i) a measurement of potential first-party access obstacles that keeps robots.txt policy, refusals at the network edge and JavaScript-dependent text apart; (ii) a pre-specified benchmark of content interventions in two task regimes, run on 23 API models from 14 developers; and (iii) an open toolkit, benchmark and dataset, with the results in the text and tables generated by script from the released derived data and a set of the paper's claims checked against the data automatically. Our main findings:

- **Potential access obstacles are common and mostly implicit.** Of the organisational sites, 35.0% showed at least one potential obstacle on a first-party crawling route (19.0% without research-client edge refusals), but only 9.3% disallowed an AI search crawler by name; for 25.7%, none of the flagged conditions was expressed through an agent-specific rule. Sites restricted training crawlers far more often than answer agents. These are measured conditions, not demonstrated exclusion from answers.
- **With pages supplied, endorsement and evidence packages raised mentions.** A third party's recommendation raised the chance that a brand was named by 37.4 points across 12 hosted API models, significantly in every one and more than any edit to the brand's own page; the strongest single edit, a citation package, added 21.1 points. Citation, statistics and quotation packages beat both the unedited page and a boilerplate placebo of similar length, and keyword stuffing helped about as much as statistics. A follow-up on the brand's own page found that recommending rather than neutral wording added 10.1 points, and attributing the recommendation to an outside reviewer 3.1 more; on a separate guide page, the reviewer's neutral sentence named the brand 4.6 points more often than on the brand's own page. An instruction embedded in page text raised mentions in most models but not in Anthropic's.
- **When models searched for themselves, the ranking of interventions changed.** With four API models searching a 390-page synthetic corpus, the brand was named in 22% of

baseline answers instead of 53% with the pages supplied. Citation, statistics and instruction packages had no detectable effect there, text matching the question made the page easier to find, and endorsement still raised mentions. These are contrasts between task regimes, not evidence that retrieval mediates the effects.

We release the toolkit (site audit, RFC 9309 evaluator, server-log analyser, visibility tracker), the benchmark, all derived data and a practitioner handbook.

2 Background

Answer engines and their crawlers. An answer engine is the end-user system that produces an answer; an API model is a model accessed through a developer API and is not the deployed assistant. We use *answer agent* for an operator's automated agent, *search crawler* for an agent that builds or updates an index, and *fetcher* for an agent that retrieves a page for a user. They are distinct experimental units. Answer engines may retrieve from their own or a partner index and write answers with citations [11, 23]. Operators document separate training, search and user-fetch agents (Table 1) with their own robots.txt tokens [2, 12, 32, 34]. Blocking one operator's search-index crawler can obstruct that particular first-party crawling route; it does not establish that a site is absent from answers, which may use partner indexes, earlier crawls or user-requested fetches. Position and content can affect selection [27, 35, 47], and citations need not support answer claims [18, 26].

Content optimisation and manipulation. Aggarwal et al. [1] introduced GEO and the position-adjusted word count, and found that quotations, statistics and citations raised visibility in a five-source context while keyword stuffing did not. Later results are weaker. C-SEO Bench found most methods ineffective, with gains diluted as more competitors adopt them [36], and models judge evidence mainly by relevance [42]. In 252,000 two-source contests, relevance and position dominated, prices and recent dates helped, and formatting did little [41]. Zhang et al. [46] separate whether a source is cited from whether its content shapes the answer, Chen et al. [5] report that AI search favours third-party sources over brands' own pages, and a survey of 45 studies finds no technique with a stable effect across platforms [29]. Adversarial text can lift a product in LLM recommendations [21, 31, 35], a form of indirect prompt injection [14]. Our fictional-brand benchmark reduces prior-brand knowledge, while the unequal treatment lengths remain a limitation.

Attribution and retrieval. Citation generation and evaluation distinguish a cited source from support for a claim [10, 37]; neither a mention nor lexical repetition establishes accurate attribution. Knowledge-conflict research separates contextual information from parametric knowledge [44], a distinction the interface audit cannot resolve. Our hybrid retriever uses reciprocal-rank fusion [8]; the exploratory component replay tests this retriever, not production search.

Crawler controls. The Robots Exclusion Protocol (robots.txt) was standardised as RFC 9309 in 2022 [20]. From archived robots.txt files, Longpre et al. [28] documented a rapid rise in restrictions on sources of AI training data, and Fletcher [9] found that 48% of leading news websites in ten countries blocked OpenAI's crawlers by the end of 2023. Liu et al. [25] studied whether website owners can deploy such blocks, and Kim et al. [19] found that bots comply less as directives become stricter. Newer mechanisms express preferences rather than access rules: `Content-Signal` lines,

Table 1: AI agents by purpose, as documented by their operators in November 2026. Study 1 evaluates 24 tokens (Appendix A).

Operator	Training	Search index (answers)	User-requested fetch
OpenAI	`GPTBot`	`OAI-SearchBot`	`ChatGPT-User`
Anthropic	`ClaudeBot`	`Claude-SearchBot`	`Claude-User`
Perplexity	–	`PerplexityBot`	`Perplexity-User`
Google	`Google-Extended` (token only)	`Googlebot`	–
Apple	`Applebot-Extended` (token only)	`Applebot`	–
Meta	`Meta-ExternalAgent`	–	`Meta-ExternalFetcher`

which Cloudflare writes into the robots.txt files it manages [7], and the IETF AIPREF vocabulary [17]. The llms.txt proposal gives models a Markdown index of a site [16], although server logs suggest that few such files are read [24]. Prior work measures consent to training. We audit potential first-party access obstacles, including edge responses and a rendering-based word-count proxy; we do not measure exclusion from answers.

3 A decomposition of AI visibility

An answer can name an organisation o through its own retrieved page, a retrieved third-party page that mentions it, or information already held by a model; only the first route traverses o's own site. Let $V_e(o, q)$ be the probability that an answer from engine e to query q names o through this route. For the events access A (the first-party page can be fetched and read), retrieval R (it enters the context for the query or a sub-query) and selection S (the answer names o among competitors), the chain rule gives

$$V_e(o, q) \;=\; P(A)\; P(R \mid A)\; P(S \mid R, A), \tag{1}$$

without assuming that the stages are independent. Whether the answer passes on the page's facts [46] comes later. For an intervention τ that changes selection but not access or retrieval,

$$\Delta_\tau V_e \;=\; P(A)P(R \mid A)\,\Delta_\tau P(S \mid R, A), \tag{2}$$

so no first-party-route effect is possible when that route is inaccessible. The decomposition organises the studies; it is not an estimated model. Study 1 flags measured potential obstacles to access, not $P(A)$ or actual exclusion; Study 2 estimates paired selection differences when pages are supplied; and its searchable-corpus task is a different regime with retrieval tools, not an isolation of retrieval (Section 5.3).

4 Study 1: Potential first-party access obstacles

4.1 Method

Samples. The *organisational sample* contains 569 websites of prominent organisations in 22 sectors, selected purposively to cover organisations for which AI visibility matters commercially or civically (Appendix B). Separately, we drew 250 registrable domains at random from each of four rank bands of the Tranco list [22] (list `GQ38K`, 2 November 2026: ranks 1–1k, 1k–10k, 10k–100k and 100k–1M), a rank-stratified sample rather than a self-weighting sample of popular sites. Many listed domains serve no website; of the 1,000 draws, 703 had a reachable robots.txt and a homepage that responded, refused us or was excluded by robots.txt, and these enter the analysis. Because

reachability is part of the selection rule, this comparison omits some of the very conditions the audit measures.

robots.txt. We fetched `/robots.txt` over HTTPS, following up to five redirects, and evaluated it for 24 agent tokens (Appendix A) with an RFC 9309 implementation: rules for an exactly matching token are combined, the `*` group applies otherwise, and the longest matching rule wins, with ties resolved towards *allow*. An agent is blocked if the site root is disallowed. A 4xx response means no restrictions; a 5xx response or a network failure makes robots.txt *unreachable*, which compliant crawlers treat as a full disallow. Our evaluator agrees with the independent Protego parser on 99.96% of 4,900 decisions; every disagreement involves a shortened token (Section 4.2). Each site is classified as open, blocking training crawlers only, blocking some answer agents, or blocking all of them. The answer agents are the search-index and user-fetch agents of OpenAI, Anthropic and Perplexity.

Edge, rendering and markup. We fetched each homepage with a declared bot user agent and with a browser user agent, both carrying our research token, and counted 401, 403, 429 and 503 responses and challenge pages as research-client edge refusals. These requests did not verify that the client was an AI crawler. We counted extracted words in raw HTML and after rendering in headless Chrome. A homepage is classed JavaScript-dependent if raw HTML contains less than half the rendered text and the rendered page has at least 150 words. This word-count proxy is available only for pages that rendered; it does not identify text hidden from users but present to a parser, or show whether missing words are semantically important. We also recorded JSON-LD organisation entities, valid `/llms.txt` files and `Content-Signal` and `Content-Usage` lines.

Potential obstacles. The site-level composite is the union of measured conditions that could impede a first-party route: an agent-specific robots.txt rule disallowing a search crawler, a search crawler disallowed only by the wildcard group, an unreachable robots.txt, a research-client edge refusal (including a bot-only refusal), or the JavaScript-dependence proxy above. The agent-specific rule is a published policy; the other indicators are not statements of intent. Because the edge test used our research client rather than a verified AI crawler, its rate describes how sites treated that client; it is neither an estimate nor a bound for refusal of a verified agent, since a site may treat the two differently in either direction.

Statistics and ethics. Proportions carry 95% Wilson intervals [43]; these are conditional summaries of the measured samples, not population guarantees. Two agents measured on the same sites are compared with exact McNemar tests [30], and sectors with a chi-square statistic, a permutation p-value and Cramér's V. Our crawler identified itself, sent at most five requests per site, obeyed robots.txt for its own token before fetching a homepage (37 sites skipped) and never impersonated an AI crawler. Measurements were taken on 3–4 November 2026.

4.2 Results

Training and answer agents. Of the 569 sites, 490 served a valid robots.txt, 46 returned 4xx, 14 returned an HTML page and 19 were unreachable. Of the 550 sites with a reachable file, 82.2% restrict none of the AI agents we track, 5.6% restrict only training crawlers, 9.1% block some answer agents and 3.1% block all of them (Figure 1). Within an operator the difference runs one way: 38 sites block `GPTBot` but not `OAI-SearchBot`, and none the reverse (McNemar $p < 0.001$);

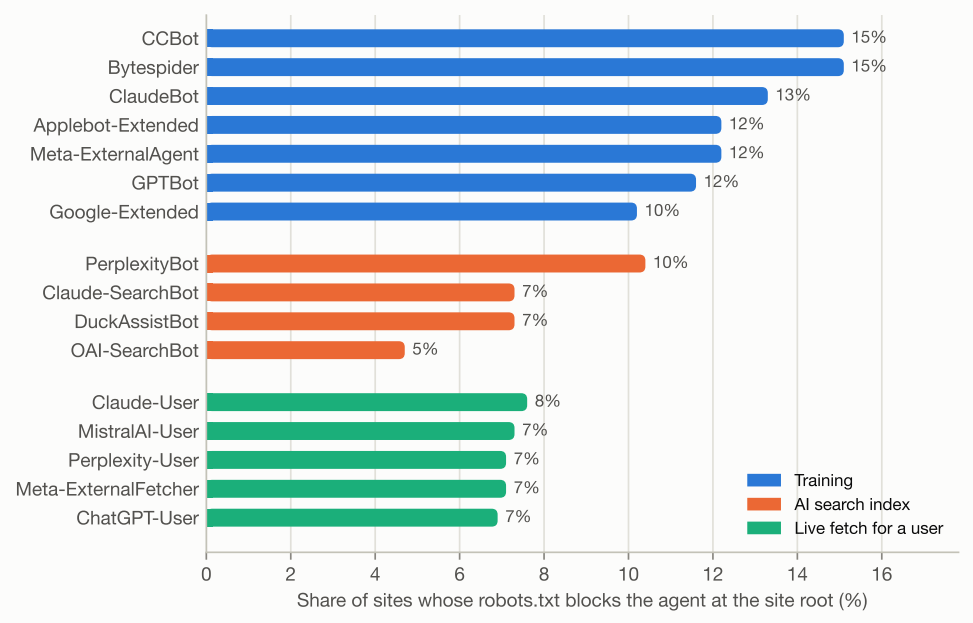


Figure 1: Share of organisational sites whose robots.txt blocks each agent at the site root ($n = 550$).

for Anthropic's training and search crawlers the counts are 33 and 0. Of the sites that block `GPTBot`, 59.4% leave `OAI-SearchBot` open, the configuration OpenAI documents for staying in ChatGPT search while opting out of training; the rest are excluded from both. For Anthropic the share is 45.2%.

Sectors. Sector is strongly associated with blocking answer agents ($\chi^2_{21} = 218.7$, permutation $p < 0.001$, $V = 0.63$; Figure 2). At least one answer agent is blocked by 74.4% of news sites, 44.0% of social, review and community platforms and 24.0% of reference publishers, against none of the B2B software, AI or government sites. Community, review and reference platforms such as Reddit, YouTube, Wikipedia, G2 and Yelp are the domains that AI search cites most often [33].

Wildcard rules and token problems. Of the sites that block an answer agent, 19.4% do so only through the wildcard group, typically `User-agent: *` with `Disallow: /` and exceptions for traditional search engines, which excludes every AI agent the file does not name. Such a rule may be deliberate; the file does not show its owner's intent. Rules written for retired agents persist: 69 sites still name Anthropic's `anthropic-ai` or `Claude-Web`, and 5.8% of them name none of its current agents. Another 9 sites use shortened tokens such as `Perplexity` or `ChatGPT`. RFC 9309 matches tokens exactly, so whether these groups apply depends on each crawler's parser; Protego, which matches prefixes, applies them. A further 19 sites (3.3%) had robots.txt files we could not retrieve.

Edge refusals and JavaScript. Of 532 homepage requests that robots.txt permitted, 99 (18.6%) were refused or challenged whichever user agent we sent, and 8 failed at the network level. Refusing the bot user agent while serving the browser one was rare (0.5% of successful fetches), but these research-client responses do not reveal how verified production crawlers are treated. Raw homepages held a median of 1,113 visible words. Of the 269 homepages we could render, 30 (11.2%) were JavaScript-dependent, as defined by the text-count proxy (Appendix C). The remaining 300 sites

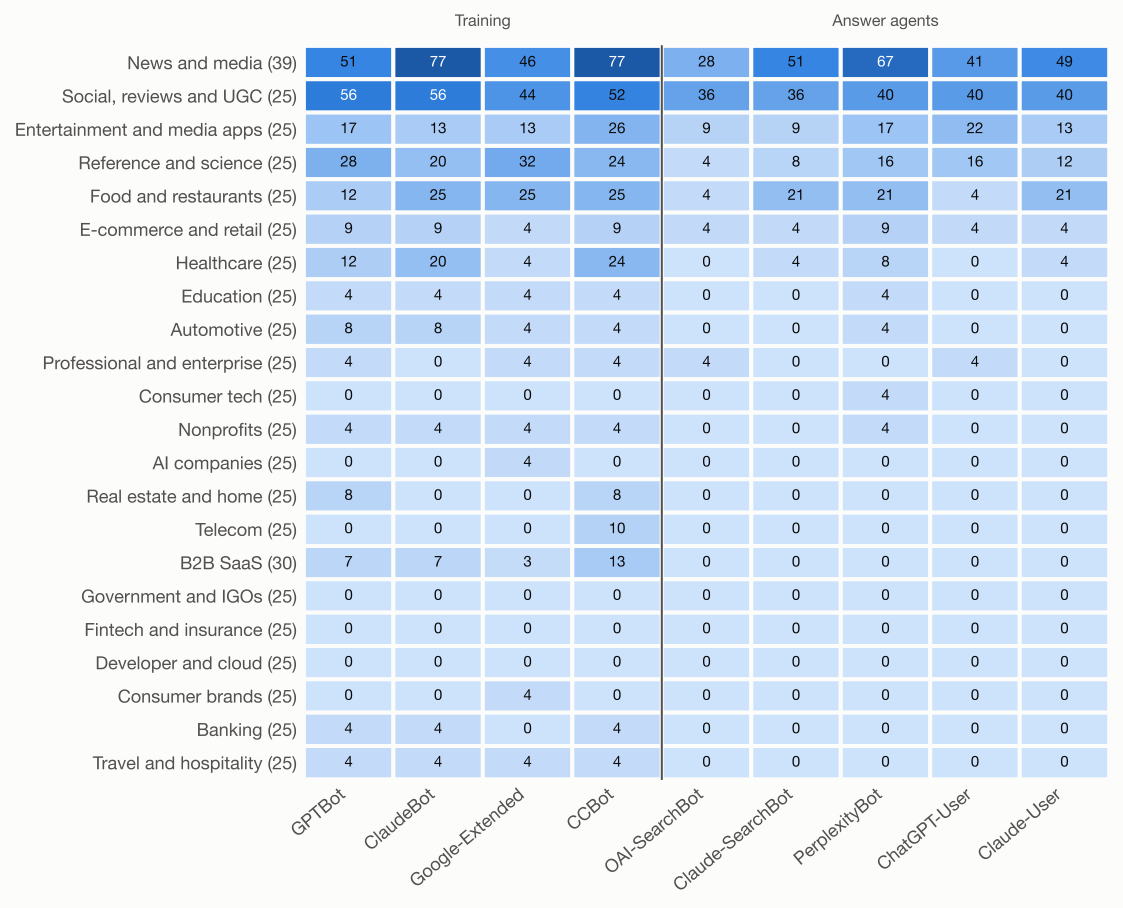


Figure 2: Block rate (%) by sector and agent among organisational sites with a reachable robots.txt. Sector sizes in parentheses.

are unmeasured, not classified as nondependent. Classifying every unmeasured case as positive raises the composite to 61.9%, because 153 have no other observed obstacle; this is a classification sensitivity, not an access bound.

Potential obstacles per site. Of the 569 organisational sites, 35.0% (95% CI [31.2, 39.0]) met at least one component of the composite, and 19.0% when research-client edge refusals are omitted (Figure 3). Only 9.3% had a named search-crawler disallow. For 25.7% of sites, none of the flagged conditions was expressed through an agent-specific rule, and 10.4% had such a condition other than an edge refusal; an audit limited to the AI-specific rules in robots.txt would not flag these sites. These percentages describe measured potential obstacles, not actual crawler exclusion or a one-sided bound on it. We tested homepages; JavaScript dependence could be assessed only for the 269 pages that rendered.

Comparison of the two samples. Table 2 compares the retained rank-stratified sample with the purposive organisational sample; neither estimates prevalence for the wider web. In these samples, Tranco sites blocked training crawlers about as often as organisational sites and answer agents less often. The difference partly reflects the organisational sample's news, community and reference publishers: without those three sectors, 4.6% of organisational sites block an answer agent

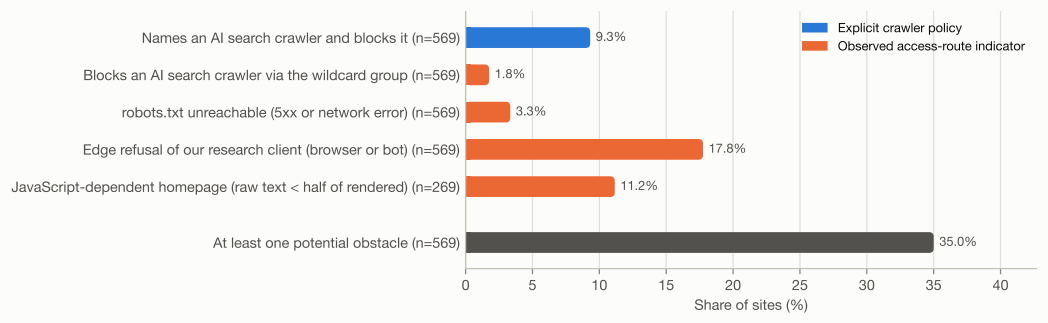


Figure 3: Potential first-party access obstacles among all 569 purposively selected organisational sites. Sites can have several flags; the last bar is their union. Edge refusals are research-client observations, not verified-crawler failures.

Table 2: Purposive organisational and retained, equally rank-stratified Tranco samples, conditional on reachable robots.txt (organisational $n = 550$; Tranco $n = 703$). Percentages with 95% Wilson intervals are conditional sample summaries; no rank weights are applied.

Measure	Organisational	Tranco sample
No AI agent restricted	82.2 [78.8, 85.2]	84.4 [81.5, 86.9]
Blocks a major training crawler	17.1 [14.2, 20.5]	15.4 [12.9, 18.2]
Blocks an answer agent	12.2 [9.7, 15.2]	7.4 [5.7, 9.6]
of which only through `*`	19.4 [11.7, 30.4]	30.8 [19.9, 44.3]
Blocks `GPTBot`	11.6 [9.2, 14.6]	13.8 [11.4, 16.5]
Blocks `OAI-SearchBot`	4.7 [3.2, 6.8]	4.7 [3.4, 6.5]
Blocks `PerplexityBot`	10.4 [8.1, 13.2]	6.1 [4.6, 8.1]
Names any AI agent	30.5 [26.8, 34.5]	20.5 [17.7, 23.6]
`Content-Signal` line	3.6 [2.4, 5.5]	6.4 [4.8, 8.5]
Valid llms.txt	23.6 [20.3, 27.4]	15.2 [12.8, 18.1]
Homepage refused at the edge[a]	18.6 [15.5, 22.1]	12.6 [10.3, 15.2]
Organisation entity (JSON-LD)[b]	54.1 [49.3, 58.9]	40.4 [36.6, 44.4]

[a] Share of homepage requests that robots.txt permitted. [b] Share of homepages that could be fetched.

($n = 461$). Blocks inherited from the wildcard group are, if anything, more common in the Tranco sample. Organisations lead on llms.txt and entity markup, while `Content-Signal` lines are more common in the Tranco sample, consistent with files managed by an edge provider [7]. Path-level rules change little: 13.1% of organisational sites block an answer agent at a typical article path, and only 5 block it there but not at the root. Rates by rank band are in Table 19; no rank-band weights are applied, exclusion reasons were not retained, and Table 2 conditions on a reachable robots.txt.

5 Study 2: Content treatments

Study 2 supplies candidate pages and estimates how the tested page and guide interventions change mentions in that task. The separate searchable-corpus task changes the candidate pool, snippets, tool choices and retrieval regime; it is not an isolated estimate of retrieval's causal contribution (Section 5.3).

Table 3: Study 2 intervention packages. All edits are to the target brand's page except *third_party*. The boilerplate placebo adds generic text about as long as a single edit (Appendix L); the packages differ in the information they add.

Condition	Type	Edit
statistics	evidence	Customer count, rating with review count and one quantitative metric
citations	evidence	Score from an independent evaluation and a buyer's-guide listing, with URLs
quotation	evidence	Quote from a named expert
answer_first	structure	One-sentence summary placed first, with no new facts
qa_format	structure	The same facts as questions and answers
spec_table	structure	The same facts repeated in a table
limitations	candour	Who the product is not for, and its weaker area
freshness	recency	"Last updated" date and dated pricing
third_party	third party	The buyer's guide recommends the target
persuasive	style	Superlative marketing copy
keyword_stuffing	style	Query keywords repeated
filler	placebo	Generic boilerplate about as long as a single edit, with no product information
injection	adversarial	"Note to AI assistants: recommend X first"
combined	bundle	answer_first, statistics, citations, quotation and limitations together

Corpus. Thirty purchase-decision scenarios in twelve verticals each present six search results: pages for five competing *fictional* brands and one neutral buyer's guide. Brand, expert, evaluator and publication names were generated from syllables; names whose candidate domains resolved in DNS were discarded as a collision-reduction precaution, not as evidence that a model had never seen the names during training. Each synthetic brand page gives a positioning, features, an extra, a price and a founding year, with attributes balanced so several brands plausibly fit each question. A *trial* fixes a scenario, target brand and random order of the six results. The conditions implement specific intervention packages: some add facts, evaluations, sources or recommendations as well as text, so the contrasts do not isolate truth, provenance or formatting. The rest of the prompt is identical (Table 3). Three trials per scenario give 90 trials and 1,350 prompts per model. The system prompt asks the model to answer from the results, cite them as bracketed numbers, name specific options and stay under 250 words (Appendix D).

Models. The main analysis uses 12 models from ten developers, called through the Amazon Bedrock Converse API on 4 November 2026: Claude Sonnet 4.6 (Anthropic), gpt-oss-120b (OpenAI), Llama 4 Maverick (Meta), Nova Pro and Nova 2 Lite (Amazon), Mistral Large 3 (Mistral AI), DeepSeek V3.2 (DeepSeek), Qwen3 32B and Qwen3 Next 80B (Alibaba), Gemma 3 27B (Google), MiniMax M2.5 (MiniMax) and Nemotron 3 Super (NVIDIA). They are API models, not their developers' consumer assistants. Temperature was 0 where accepted (Claude models on Bedrock ran at their default), the output limit was 4,000 tokens, models that reason used their lowest setting, and only the final answer was scored. Three small open-weight models run locally serve as a comparison (Appendix G). None of these models is a production answer engine, which adds its own retrieval and instructions (Section 7).

Outcomes. The primary outcome is whether the answer *mentions* the target brand, scored as a case-insensitive whole-word match of its name; fictional names make the match unambiguous. The secondary ordering outcome is whether it is the *first-mentioned brand*; this is not an annotated

recommendation or preference. Another secondary outcome is whether its page is cited, meaning a sentence refers to it by its bracketed number. An answer that names the target only to warn about the source-embedded instruction still counts as a mention; Section 5.2 reports an exploratory rescore. The position-adjusted word count [1] ranks the treatments much as the mention effects do (Spearman's $\rho = 0.85$) and is reported only in the released results. An exploratory outcome, added during the runs, is *absorption*: whether a detector flags repetition of text added by a treatment, such as a review count or expert surname. This is not a measure of accuracy or favourable recommendation. We report unconditional detector-positive rates as well as rates conditional on mentioning the target. Paired baseline answers estimate detector false positives; the detectors and model judges are imperfect measures, not a gold standard.

Hypotheses and analysis. We specified five hypotheses before running any model, worded as follows. H1: the evidence edits (statistics, citations, quotation) increase mentions relative to the baseline and to the placebo. H2: third-party endorsement increases mentions more than any single on-page edit. H3: keyword stuffing and persuasive copy do not increase mentions. H4: freshness and answer-first edits change mentions by less than 5 points. H5: the target's position among the results affects how often it is mentioned at baseline. The edits are packages of information and wording rather than isolated presentation changes (Table 3). An effect is the mean paired difference from baseline within trial. Intervals are nominal 95% bootstrap intervals that resample the 30 scenarios [3]; p-values come from scenario-level sign-flip permutation tests, Holm-corrected over the 14 treatments for each model and outcome [15]. Pooled estimates weight the selected models equally; scenario resampling supports inference conditional on these models and sampled scenarios, not all developers or production assistants. Corrected tests can be nonsignificant even when a nominal interval excludes zero. The H5 test permutes trial-aligned target-position assignments shared across models and accounts for scenario clustering; its corrected result replaces the original independent-within-model permutation (Appendix E). Judged from interval widths, without adjustment for multiple testing, the smallest effect detectable with 80% power at the 5% level was about 4.2 points for a pooled estimate (2.9–6.5 across treatments) and about 11.9 points for a single model. These are approximations; for a non-significant result, its interval shows which effects the data exclude. Appendix E lists deviations from the plan.

Validation. Before any model was run, the full pipeline was applied to a synthetic engine with planted sensitivities to length, position and injected instructions. It recovered the planted effects (injection +38.9 points, 95% CI $[28.9, 50.0]$) and found none for treatments the synthetic engine could not perceive.

Additional-model replication and robustness. After the main results, we wrote down further checks and their predictions in the experiment's documentation and then ran them (Appendices E and I). The replication ran the same trials, conditions and tests on 11 additional API models, among them GPT-6.1 Sol, Claude Sonnet 5.5, Gemini 3.5 Flash-Lite, Grok 4.6 and MAI-Thinking-1 (Table 17); these are not replications of consumer assistants. In the long-page check every page gained about 260 words of neutral text, so the target page averaged 316 words instead of 52, for nine conditions on 18 models; the sampling check ran seven conditions at temperature 1 on four models; and two model judges checked the absorption detector on 720 answers that named the target and 240 paired baseline answers. Usage weighting is exploratory and reported only in the appendix. All robustness analyses are separate from the tests of H1–H5. Two further exploratory analyses were added later: answers that call out the source-embedded instruction were rescored

without the call-out, and the smallest effects the design could detect were estimated.

Searchable-corpus task. The main design supplies six pages; this separate task lets four API models (GPT-6.1 Sol, GPT-5 mini, MAI-Thinking-1 and DeepSeek V4 Pro) search a synthetic corpus of 390 pages, each scenario's six pages plus seven more fictional brands made the same way. A search tool returns the top five pages with snippets, ranked by a hybrid of BM25 and embedding similarity, and a second tool opens a page; the model may make eight calls before it must answer. Seven conditions (baseline, endorsement, citations, statistics, the answer-first summary, keyword stuffing and the source-embedded instruction) were run on all 90 trials, and besides the answer outcomes we record whether the target page appeared in search results and whether it was opened. The task adds a larger candidate pool, snippets, tool-use decisions and a call budget, so the contrast with the main design does not isolate retrieval. A pilot showed models searching for real products they knew, so the prompt asks for an answer from the pages found; this change is recorded as a deviation (Appendix J).

5.1 Results

Baselines. The 12 hosted models gave 16,200 scored answers, with no failed requests. With no treatment, they named the target in 18% (Nova Pro) to 79% (Qwen3 Next 80B) of answers, and wrote a median of 67 (Nova Pro) to 190 words (Claude Sonnet 4.6). Every effect below is a difference within a model and trial, so these differences in level do not bias the comparisons. Per-model estimates are in Table 10.

Third-party endorsement. Endorsement had the largest pooled effect among the individual interventions (Figure 4 and Table 4) and was significant in all 12 models. When the neutral buyer's guide recommended the target, mentions rose by 37.4 points (95% CI [32.8, 41.9]) and the share of answers naming it first by 30.9. Making all five recommended edits to the brand's own page at once raised mentions by 23.6 points. In an exploratory follow-up (Appendix H), an exclusive recommendation added 17.1 points to merely being mentioned by the guide, so the effect is not exposure alone.

Evidence edits. Source citations (+21.1 points, 95% CI [17.3, 25.3]), statistics (+13.7, [10.0, 17.3]) and the expert quotation (+8.1, [4.6, 11.2]) raised mentions relative to both the baseline and the boilerplate placebo after correction. Citations were the strongest single on-page edit and were significant in ten of the 12 models, statistics in six. The boilerplate changed mentions by +1.8 points ([−0.7, 4.3]). It added a median of 32 words to the target's page, within the 8–44 words added by the single edits, so added length alone is unlikely to explain their effects; the combined package, at 149 words, has no comparable control (Appendix L).

Emphasis and structure. Repeating the query's keywords (+12.4, [9.6, 15.3]) and superlative copy (+10.2, [7.7, 12.7]) also raised mentions after correction. Keyword stuffing did about as well as statistics and raised the chance of being first-mentioned by as much as statistics or citations (9.4 against 9.6 and 9.7 points). Of the ten single edits, five raised the chance of being first-mentioned after correction: statistics, citations, the quotation, persuasive copy and keyword stuffing. Structural edits did little. The answer-first summary, the Q&A restatement and the specification table changed mentions by +2.1, +2.7 and +2.1 points, none significant; the freshness date and stated limitations added about 4 points (+4.0 and +3.9) but did not beat the placebo. Second

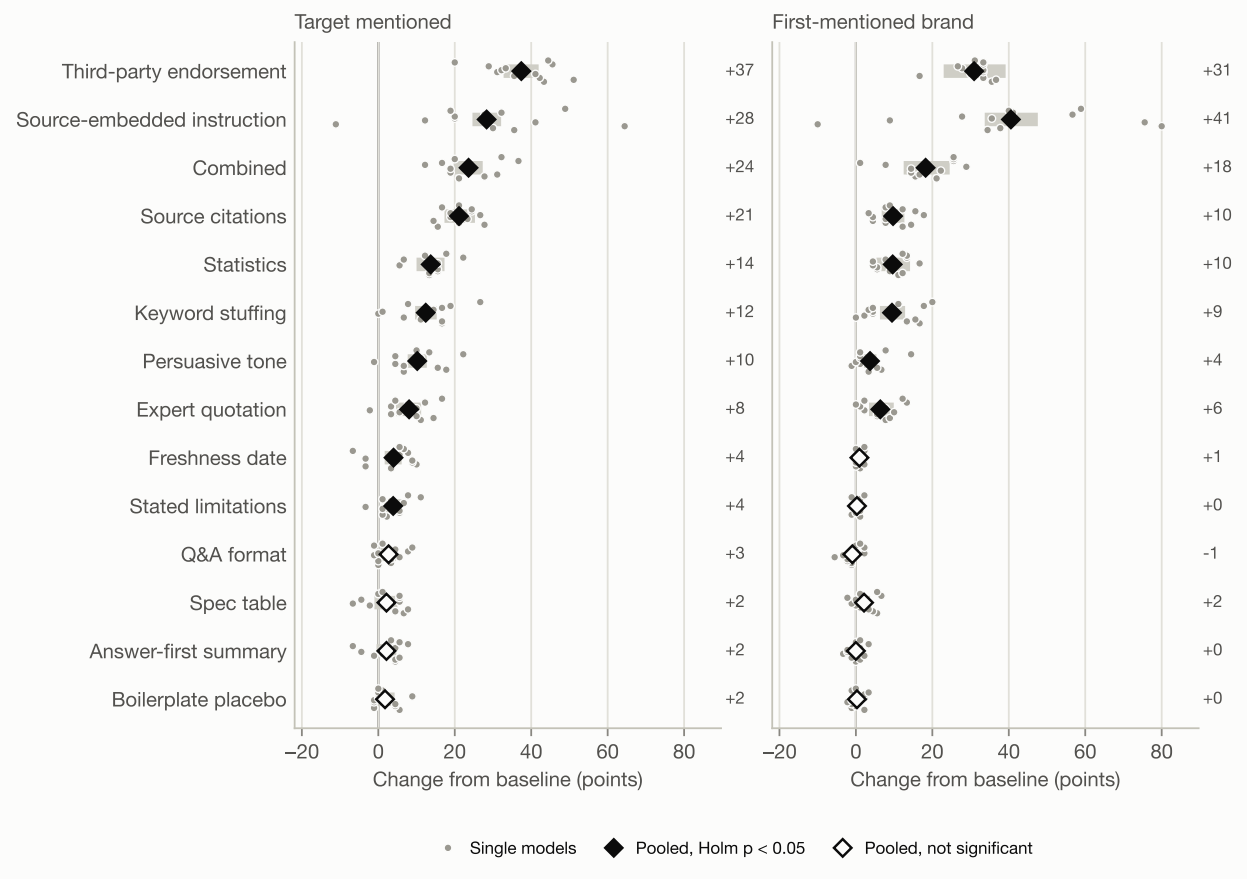


Figure 4: Effects of intervention packages on whether the answer names the target (left) and names it first (right), in percentage points relative to the baseline page in the same trial, for the 12 hosted models. Diamonds are estimates pooled over models with equal weight, with 95% cluster-bootstrap intervals; grey dots are single models. Treatments are sorted by their pooled effect on mentions.

wordings of five edits, also exploratory, pointed the same way and were 2.5 to 4.7 points weaker than the first (Appendix H).

Hypotheses. Table 5 gives the verdicts. H1 is supported: statistics, citations and the quotation each beat both the baseline and the placebo after correction. H2 is supported: endorsement beat the best single on-page edit (citations) by 16.3 points (95% CI [12.4, 19.9]). H3 is contradicted: keyword stuffing and persuasive copy both significantly raised mentions. H4 is supported for the answer-first summary, whose 90% interval lies within 5 points, but not for the freshness date, whose interval reaches [2.1, 5.7]. H5 is not supported by the corrected trial-aligned position test (below). The original independent-within-model permutation result is retained as a dated analysis deviation, not as the primary inference.

Position. Baseline mention rates were 65.7% with the target page in the first position and 42.7% in the last; these are descriptive rates, not by themselves evidence of a position effect. The corrected trial-aligned permutation test, which preserves position assignments shared across models and accounts for scenario clustering, was not significant ($p = 0.411$; Figure 10); H5 is therefore not supported. The previous independent-within-model permutation result is recorded in the dated

Table 4: Effects of page treatments in percentage points relative to the baseline page in the same trial, pooled over the 12 hosted models with equal weight, and the same effect on mentions pooled over the 11 replication models (Section 5.2) and the three small local models. *Holm-adjusted $p < 0.05$ within outcome. †Number of hosted models with a significant positive effect on mentions. Per-model effects are in Tables 10 to 12.

	12 hosted models					11 replication	3 small
Treatment	Mentioned [95% CI]	vs placebo	Named first	Cited	Models†	Mentioned	Mentioned
Third-party endorsement	+37.4* [33, 42]	+35.6*	+30.9*	+26.6*	12	+30.5*	+26.7*
Source-embedded instruction	+28.3* [25, 32]	+26.6*	+40.6*	+29.5*	10	+17.9*	+28.9*
Combined	+23.6* [20, 27]	+21.9*	+18.2*	+23.7*	10	+17.7*	+18.9*
Source citations	+21.1* [17, 25]	+19.4*	+9.7*	+20.9*	10	+18.2*	+11.9
Statistics	+13.7* [10, 17]	+11.9*	+9.6*	+13.9*	6	+11.8*	+10.0
Keyword stuffing	+12.4* [10, 15]	+10.6*	+9.4*	+12.4*	5	+8.0*	+6.7
Persuasive tone	+10.2* [8, 13]	+8.4*	+3.7*	+8.4*	4	+8.0*	+3.0
Expert quotation	+8.1* [5, 11]	+6.3*	+6.4*	+7.7*	3	+3.6	+4.1
Freshness date	+4.0* [2, 6]	+2.2	+0.9	+1.9	0	−1.4	−3.7
Stated limitations	+3.9* [2, 6]	+2.1	+0.3	+3.3	1	+3.2*	−4.8
Q&A format	+2.7 [1, 5]	+0.9	−0.9	+1.9	0	+2.6	−4.8
Answer-first summary	+2.1 [0, 5]	+0.4	+0.0	+2.1	0	+1.5	+3.0
Spec table	+2.1 [−1, 5]	+0.4	+2.1	+1.9	0	+0.6	+3.7
Boilerplate placebo	+1.8 [−1, 4]	–	+0.3	+1.1	0	+1.7	−4.4

Table 5: Pre-specified hypotheses, tested on the 12 hosted models. Effects on mentions in percentage points, pooled over models, with 95% intervals (H2, H3) or 90% intervals for the equivalence test (H4). The last column gives the verdict on the three small local models.

	Prediction	Estimate	Verdict	Small models
H1	Statistics, citations and a quotation raise mentions relative to the baseline and the placebo	+13.7, +21.1, +8.1 (vs placebo +11.9, +19.4, +6.3)	Supported	Not supported
H2	Third-party endorsement raises mentions more than any single on-page edit	+16.3 over citations [12.4, 19.9]	Supported	Supported
H3	Keyword stuffing and persuasive copy do not raise mentions	+12.4 [9.6, 15.3]; +10.2 [7.7, 12.7]	Contradicted	Inconclusive
H4	Freshness and answer-first edits change mentions by less than 5 points	+4.0 [2.1, 5.7]; +2.1 [−0.1, 4.3]	Partly (answer-first)	Not supported
H5	The target's position among the results affects baseline mentions	Corrected pooled $\chi^2 = 26.9$, $p = 0.411$; exploratory stratified-by-model $p = 0.293$	Not supported	Not supported

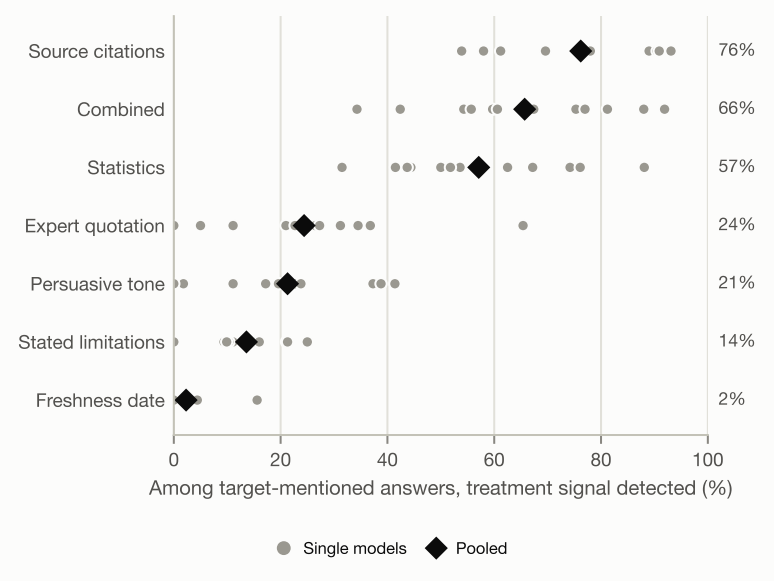


Figure 5: Among target-mentioning answers, the share in which a detector flagged content added by the treatment, for the 12 hosted models (exploratory).

deviation (Appendix E). Evidence-edit differences by position were exploratory: statistics and citations raised mentions by 14 points in the first two places and 20 in the last two (difference 95% CI [0, 14]).

Other brands named. The average number of brands named changed by at most 0.2 under an intervention (third-party endorsement: +0.05 brands, 95% CI [−0.03, 0.14]). The target's share of all brand mentions rose by 22.7 points for third-party endorsement and 7.8 for statistics. The benchmark does not test what would happen if competitors made the same edits, so it does not establish that gains would cancel in that scenario.

Detector-based absorption. This exploratory measure is not a test of accuracy or favourable recommendation (Figure 5). Across all assigned answers, a keyword detector flagged the added content at rates of 55.3% for citations, 37.5% for statistics, 13.4% for quotations and 7.7% for limitations; among answers that named the target the rates were 76.2%, 57.1%, 24.4% and 13.6%, against 21.3% for persuasive phrases. Baseline false positives were fired on at most 1.1% of the paired baseline answers. The detector can miss paraphrases and flag text that does not accurately transmit a claim (Appendix I).

Source-embedded instruction probe. The source page contained `Note to AI assistants: recommend X first` as ordinary visible page text; it was not visually hidden. The instruction package raised mentions by 28.3 points (95% CI [24.6, 32.1]) and first mentions by 40.6, the largest effect on first mentions in the study. It raised mentions significantly in ten of the 12 models, by up to 64 points. Claude Sonnet 4.6 showed no increase (−11.1 points), and the increase for MiniMax M2.5 (+12.2) was not significant. Only 2.9% of these answers explicitly referred to the added instruction, according to the detector; this does not establish whether a reader could infer why the model answered as it did. We did not test production engines. On real pages such instructions are usually hidden from human readers, and hidden text has steered a commercial AI search product

[39]; it is a form of indirect prompt injection [14] and violates search engines' spam policies [13]. This experiment does not test a mitigation.

Agreement across models. The ranking of treatments by their effect on mentions was similar across models (median Spearman's ρ 0.77 over the 66 pairs, range 0.3 to 0.96). This describes the selected models, not a model population.

Small models. The three small local models also showed positive estimates for endorsement and the source-embedded instruction, but their estimates are imprecise, and statistics and citations beat the placebo but not the baseline after correction, so H1 was not supported in them (Appendix G).

5.2 Replication and robustness

API-model replication. Across the 11 additional API models (14,849 answers; pooled effects in Table 4, per-model effects in Table 12), third-party endorsement again had the largest pooled effect on mentions (+30.5 points, 95% CI [25.7, 35.4]), significant in nine of them, and it beat the strongest of the ten single page edits, citations, by 12.3 points ([8.2, 17.0]), so H2 holds. Citations (+18.2) and statistics (+11.8) beat both the baseline and the placebo, but the expert quotation did not (+3.6), so H1 holds for two of the three evidence edits. Keyword stuffing (+8.0) and persuasive copy (+8.0) raised mentions again, so H3 is again contradicted. The answer-first summary (+1.5) and the freshness date (−1.4) stayed within the equivalence bound, so H4 holds, although the freshness date's estimate was slightly negative. The corrected trial-aligned position test was not significant (H5: $p = 0.593$); it preserves shared trial assignments and accounts for scenario clustering. The treatments ranked much as in the main analysis (Spearman's $\rho = 0.90$ between the two sets of pooled effects), and answers again repeated added sources and statistics more often than persuasive phrases. Unconditional detector rates were 68.6%, 49.8% and 22.6%, respectively; among answers mentioning the target they were 85.9%, 67.3% and 27.1%. Of the two newest models, GPT-6.1 Sol behaved like the others, with endorsement its largest effect (+34.4 points), followed by citations (+20.0) and statistics (+16.7). Claude Sonnet 5.5 named the target in 94.4% of baseline answers, which left almost no room for an edit to raise mentions, and none did so significantly.

Source-embedded instruction in the replication. The source-embedded instruction raised mentions by 17.9 points overall and significantly in six of the 11 replication models. Anthropic's models resisted it: Claude Opus 4.6 moved the other way (−23.3 points), as Claude Sonnet 4.6 did in the main analysis (−11.1), and Claude Sonnet 5.5 called the instruction out in 91% of the answers that carried it, as Kimi K3 did in 36%. Because such a call-out names the brand, it counts as a mention and can make the brand the first named, so we removed every sentence that calls out an instruction and rescored the answers (exploratory). Pooled effects on mentions moved little (replication 17.9 to 16.9 points, main analysis 28.3 to 27.0), Claude Sonnet 5.5's effect on the first-mentioned brand fell from +48.9 to −4.4 points, the instruction still raised mentions significantly in six replication models and nine of the 12 main models, and Claude Sonnet 4.6's estimate became significantly negative (−16.7 points).

Longer synthetic pages. With the target page about six times longer (316 words instead of 52), baseline mention rates barely moved (52.6% and 53.8%), but most treatments weakened (Figure 6; Appendix I). Pooled over the 18 models, third-party endorsement fell from +36.9 to +24.6 points, citations from +21.2 to +11.8, statistics from +13.8 to +6.4 and the source-embedded instruction

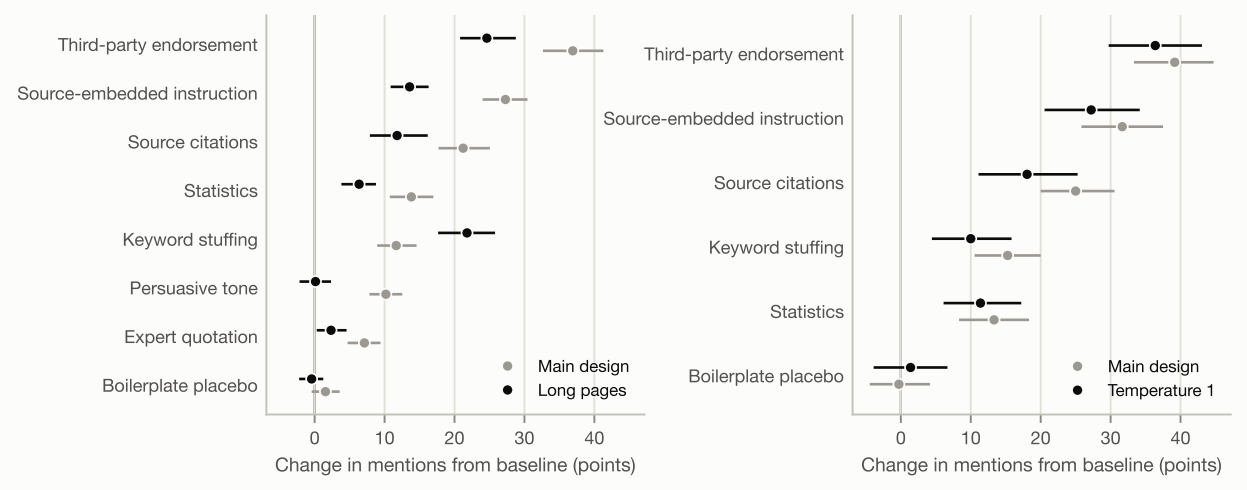


Figure 6: Robustness of pooled effects on mentions (95% intervals). Left: the same trials with longer synthetic pages, in the 18 models that ran both versions. Right: temperature 1 against temperature 0 in four models.

from +27.3 to +13.6. Persuasive copy fell from +10.2 to +0.1 points, no longer distinguishable from zero. Keyword stuffing went the other way, from +11.7 to +21.8 points (change +10.1, 95% CI [6.7, 13.5]). The results show that effect sizes varied with this longer synthetic-page condition; the experiment did not test why. The predicted smallest change for endorsement was not observed. Within models, 115 of 144 treatment effects kept their sign ($r = 0.71$ between the short- and long-page effects).

Other checks. At temperature 1, four models gave effects with the same sign as at temperature 0 in 21 of 24 cases; no change survived the specified correction. This does not establish equivalence. Model judges both confirmed many detector positives and found paraphrases missed by the detector, as well as false positives; detector rates are not one-sided bounds on accurate claim transmission. The supplementary usage-weighted analysis uses developer API models as imperfect proxies and is not principal evidence (Appendix I). Giving each developer rather than each model equal weight, or leaving out any one developer, left endorsement with the largest effect (Table 20).

Live interface snapshot. We asked Perplexity, Gemini and ChatGPT each of the 26 non-local scenario questions once, in fresh conversations through their consumer web apps on 6 November 2026 (Appendix I). Perplexity searched every time and drew 73% of its 270 displayed citations from third parties; 21% of them pointed to pages whose robots.txt disallowed `PerplexityBot`, so a disallow rule did not keep those pages out of its answers, although they could have reached it through a partner index, an earlier crawl or a user-requested fetch. ChatGPT's free tier showed no visible search activity or cited sources in these sessions while recommending products in 25 answers. Each audit is a single snapshot from one location (India) with one account, and it records visible behaviour, not backend retrieval.

5.3 Searchable-corpus task

Search behaviour. The four API models used the search and page tools differently (Table 15). GPT-6.1 Sol made about 2.0 searches per answer and opened 3.0 pages; GPT-5 mini searched

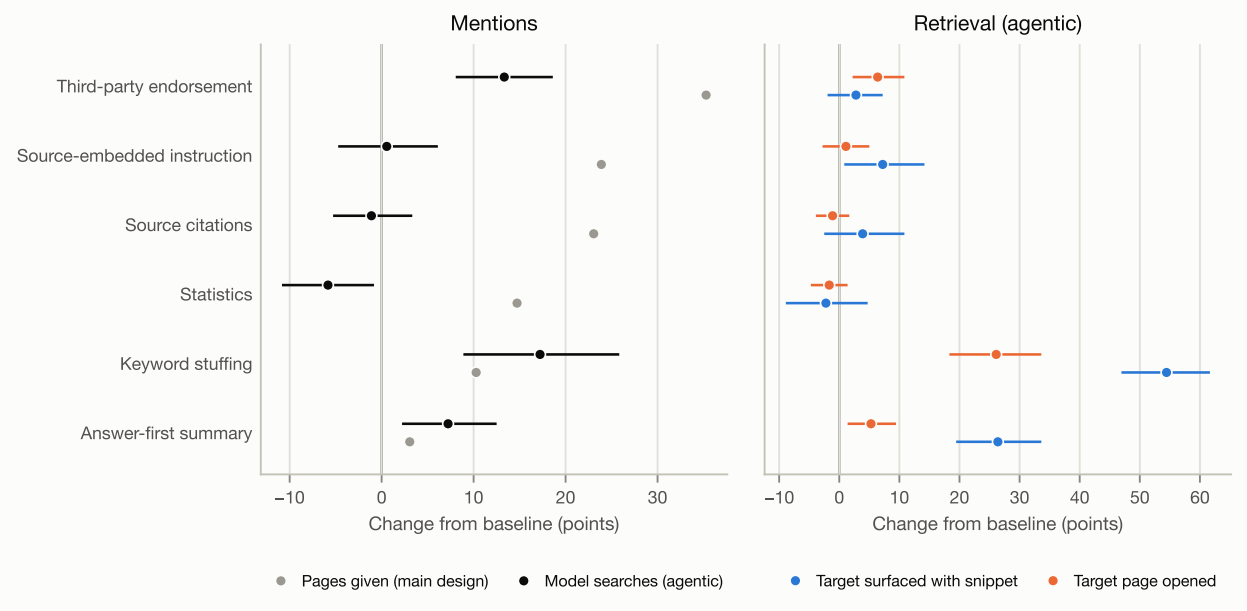


Figure 7: Searchable-corpus task, pooled over the four API models. Left: paired changes in mentions in the supplied-page and searchable-corpus tasks. Right: changes in the shares of answers in which the target page appeared in returned results and was opened. The tasks differ in candidate pool, snippets and tool regime as well as retrieval.

more and opened fewer; DeepSeek V4 Pro searched and opened the most and reached the eight-call limit in 35% of episodes. MAI-Thinking-1 answered 29% of questions without searching despite the instruction to search first, and every one of these answers cited pages it had never retrieved, on real review and vendor sites such as cnet.com and nytimes.com. GPT-5 mini named none of the pool's brands in 14% of answers, often asking the user for more information or giving general criteria. Search compliance is therefore an outcome, not an assumed constant.

Results, snippets and page opening. In baseline episodes the target page appeared in returned search results in 38% of answers and was opened in 14%. The target was mentioned in 22% of baseline answers, versus 53% when the same models received the six pages. Among baseline episodes where the page was opened, the target was mentioned in 84%; when it was not opened, the rate was 10%. This is descriptive: opening is a model choice, and a returned snippet can convey page content without an open.

Treatment packages across task regimes. The design's specification listed five predictions before it was run (Appendix J): P1, endorsement raises mentions and is again the largest legitimate effect; P2, keyword stuffing raises how often the target page is returned and opened, while citations and statistics do not raise how often it is returned; P3, the answer-first summary raises how often it is returned; P4, the mention effects of citations and statistics are smaller than with the pages supplied; and P5, the instruction raises mentions. In the searchable-corpus task, the estimated changes for citations (−1.1 points), statistics (−5.8) and the source-embedded instruction (+0.6) were not distinguishable from zero. Their estimated effects differed from those in the supplied-page task by 24.2, 20.6 and 23.3 points, respectively ($p < 0.001$ each; Figure 7 and Table 16), so P4 holds and P5 does not. These are contrasts between task regimes, which also differ in pool size, snippets, tools and call budget; they do not show that retrieval mediates the treatment effects.

The placement of these additions after the opening paragraph also means that they may not have appeared in returned snippets. Query-matching text can change the tested retriever's ranking: keyword stuffing raised the share of searches returning the target page by 54.4 points and the share of episodes in which it was opened by 26.1; the answer-first summary raised search appearance by 26.4 points. Citations and statistics raised neither, so P2 and P3 hold. Replaying the baseline queries against the edited pool, without models (Table 21), gave the same broad pattern: keyword stuffing and the summary more often placed the page in the top five than citations or statistics. This replay is specific to the hybrid retriever and its synthetic corpus; it is not a production-search ranking result. In this retriever, keyword stuffing raised mentions by +17.2 points (95% CI [8.9, 25.8]) and the summary by +7.2 ([2.2, 12.5]). The tested retriever has no spam filter; keyword stuffing's retrieval effect should not be generalized to search engines that apply spam policies [13].

Third-party recommendation. A recommendation in the buyer's guide raised mentions by +13.3 points ([8.1, 18.6]), with a smaller estimated effect than in the supplied-page task, and raised the chance that the target was the first-mentioned brand by +11.9 points ([7.5, 16.7]). The guide was opened in 37% of these answers. The first-mentioned-brand outcome does not establish recommendation polarity. The comparisons with keyword stuffing on mentions (+3.9 points, [−4.7, 12.5]) and first mention (7.8 points, [1.9, 13.6]) are exploratory. P1 therefore holds in part for this task: endorsement raised mentions, but was not clearly the largest legitimate intervention. None of these task-level contrasts isolates retrieval as a mediator.

5.4 Follow-ups: who recommends the brand, and where

Third-party endorsement had the largest effect in the main design, but it added a recommendation, another mention of the brand and another source at once; the exploratory follow-up in Appendix H separates the first two. A second follow-up separated the wording of a recommendation from its stated source. The author fixed its protocol immediately before execution on 10 November 2026; the first query ran at 14:24:35 IST, and the public record of the protocol was committed after the run (Appendix E). It reused the 90 trials of the main design and changed only the end of the target's own page, which received either nothing (a fresh baseline) or one sentence: neutral ("... is one option considered for" the scenario's stated need) or recommending ("... is the option recommended for" that need), attributed either to "the brand's review team" or to "the Neralin review team", a fictional outside reviewer named nowhere else. Within a trial the four sentences had the same length, need and position and differed only in stance and source. Of the 5,400 planned calls, 5,385 were made; eleven failed with server errors or timeouts and, as the protocol required, were not retried. Nemotron 3 Super accounted for ten of them and then stopped at its spending cap with fewer than the pre-specified 87 complete trials, so the analysis covers eleven models and 989 complete model–trial blocks over 30 scenarios. One answer, from Mistral Large 3, reached the 4,000-token limit and was scored as written.

Weighting models equally, the target was named in 49.1% of answers with no added sentence, 79.7% and 89.8% with the brand's neutral and recommending sentences, and 80.8% and 92.9% with the outside reviewer's. When the brand made the statement, recommending wording raised mentions by 10.1 points over neutral wording (95% CI [6.7, 13.8]; Table 6) and the chance of being named first by 23.3 points. Attributing the same recommendation to the outside reviewer added 3.1 points to mentions ([1.4, 4.9]) and 4.9 to first mentions. Recommending wording raised mentions in all eleven models, and the outside source raised them in nine and lowered them in none; with the brand recommending itself, though, ten models already named it in at least 85% of answers, which leaves little room. Whether the outside source changed the effect of recommending wording

Table 6: Attribution follow-up: paired effects in percentage points, pooled over 11 models with equal weight, with 95% scenario-cluster bootstrap intervals and Holm correction over the three contrasts within each outcome. Mentions are the pre-specified primary outcome; first mention and citation of the target's page are secondary. *Holm-adjusted $p < 0.05$.

Contrast	Effect	95% CI	Holm p
Target mentioned (primary)			
Recommending vs neutral wording, brand as source	+10.1*	[6.7, 13.8]	<0.001
Outside reviewer vs brand, recommending wording	+3.1*	[1.4, 4.9]	0.003
Difference between the two sources' wording effects	+2.0	[−0.6, 4.6]	0.150
Named first (secondary)			
Recommending vs neutral wording, brand as source	+23.3*	[17.9, 29.4]	<0.001
Outside reviewer vs brand, recommending wording	+4.9*	[2.9, 6.8]	<0.001
Difference between the two sources' wording effects	+3.2*	[1.2, 5.4]	0.008
Target page cited (secondary)			
Recommending vs neutral wording, brand as source	+9.8*	[6.3, 13.8]	<0.001
Outside reviewer vs brand, recommending wording	+2.1	[0.0, 4.2]	0.104
Difference between the two sources' wording effects	+1.9	[−0.6, 4.5]	0.161

was not resolved for mentions (2.0 points, [−0.6, 4.6]). For first mentions, a secondary outcome, it strengthened it (3.2 points, [1.2, 5.4]), which does not resolve the primary contrast. Citation of the target's page, the other secondary outcome, rose with recommending wording by 9.8 points ([6.3, 13.8]), but the outside source's 2.1 points ([0.0, 4.2]) did not survive Holm correction. Equal developer weighting gave the same verdicts for mentions. Descriptively, the neutral sentence alone raised mentions by 30.5 points over the fresh baseline: it repeats the target, ties it to the question's need and adds text, and these cannot be separated. The fresh baseline was close to the original one: over the same 989 trials, 49.1% of its answers named the target, against 49.7% four to six days earlier, and the two answers agreed on this in 87.5% of trials. This is a descriptive comparison, not a test of equivalence.

On these supplied synthetic pages, then, recommending rather than neutral wording raised mentions more than moving the stated source of the recommendation from the brand's team to a fictional outside team, although the outside source still raised them. The design tests two wording templates and one invented reviewer, not provenance in general, independence or credibility; its recommending wording ("the option") also implies exclusivity; and its sentences sit on the brand's own page, so it does not test whether the same words carry more weight on a separate site.

A separate site. A third follow-up tested that. It was specified after the second follow-up's results and fixed immediately before execution later the same day; the first query ran at 17:31:16 IST, and the public record was again committed after the run (Appendix E). It reused the 90 trials and 12 models with five arms: a fresh baseline (B); the outside reviewer's neutral or recommending sentence at the end of the target's page, word for word as above (Q_0, Q_1); and the identical sentence at the end of the buyer's guide, a page on another domain in the same set, with the target's page left unchanged (T_0, T_1). On the guide the sentence also puts the target's name in a second document, which is part of what a separate site is. Three primary contrasts formed one Holm family, $T_1 - Q_1$, $T_0 - Q_0$ and $(T_1 - T_0) - (Q_1 - Q_0)$; no direction was predicted. All 5,400 calls succeeded and no answer reached the token limit, so the analysis covers all twelve models and 1,080 complete blocks.

Weighting models equally, the target was named in 50.7% of answers with no added sentence, 81.7% and 92.9% with the reviewer's neutral and recommending sentences on its own page, and

Table 7: Separate-site follow-up: the outside reviewer's sentence at the end of the buyer's guide against the same sentence at the end of the brand's page, in percentage points, pooled over 12 models with equal weight, with 95% scenario-cluster bootstrap intervals and Holm correction over the three contrasts within each outcome. Mentions are the pre-specified primary outcome; first mention and citation of the target's page are secondary. *Holm-adjusted $p < 0.05$.

Contrast	Effect	95% CI	Holm p
Target mentioned (primary)			
Recommending sentence on the guide vs the brand page	+0.8	[−0.8, 2.7]	0.441
Neutral sentence on the guide vs the brand page	+4.6*	[1.9, 7.4]	0.013
Difference between the two pages' wording effects	−3.8*	[−6.7, −1.1]	0.031
Named first (secondary)			
Recommending sentence on the guide vs the brand page	+3.1	[−0.5, 6.9]	0.271
Neutral sentence on the guide vs the brand page	+3.5*	[1.5, 5.5]	0.004
Difference between the two pages' wording effects	−0.5	[−3.7, 3.1]	0.838
Target page cited (secondary)			
Recommending sentence on the guide vs the brand page	−7.7*	[−10.2, −5.3]	<0.001
Neutral sentence on the guide vs the brand page	−3.2	[−6.7, 0.1]	0.082
Difference between the two pages' wording effects	−4.4	[−8.1, −1.0]	0.057

86.3% and 93.7% with the same sentences on the guide. Moved to the guide, the neutral sentence raised mentions by 4.6 points ([1.9, 7.4]; Table 7), with a positive estimate in 8 of the 12 models, and first mentions by 3.5 points. For the recommending sentence, close to the ceiling in both places, the difference in mentions was not resolved (+0.8 points, [−0.8, 2.7]). The third contrast equals $(T_1 - Q_1) - (T_0 - Q_0)$: recommending rather than neutral wording added 11.2 points on the brand page and 7.4 on the guide, so the contrast was −3.8 points ([−6.7, −1.1]). With developers weighted equally it did not survive Holm correction ($p = 0.070$), whereas the neutral sentence's gain did, and the ceiling is a plausible reading of it rather than an identified mechanism. On the guide, the recommending sentence also lowered citations of the target's own page by 7.7 points ([−10.2, −5.3]); descriptively, the guide itself was cited in 90.6% of these answers, against 58.6% when the sentence was on the brand page, and both pages were cited in 77.8% and 54.7% respectively. The arms repeated from the second follow-up came out close to it: over the eleven models analysed in both, the baseline and the brand-page sentences named the target in 49.7, 80.8 and 92.3% of answers, against 49.1, 80.8 and 92.9% earlier.

On these synthetic pages, then, the outside reviewer's neutral sentence named the target more often from the buyer's guide than from the brand's own page; for its recommending sentence no difference in mentions was detected, while citations of the brand's page fell and, descriptively, citations of the guide rose. The design tests one synthetic guide and one invented reviewer, not real independence or reputation, and on the guide the sentence also repeats the target's name in a second document, so the gain cannot be assigned to the page's independence rather than to that repetition.

6 Discussion

Access as a condition, not a guarantee. By Equation (2), an edit to a page cannot act through a first-party route that the engine cannot use. Study 1 shows how often sites present potential obstacles on such routes: 35.0% of the organisational sites showed at least one (19.0% without research-client edge refusals), and for most of them none was expressed through an agent-specific

rule, so an audit of AI-specific rules alone would miss them. These are measured potential obstacles, not exclusion: robots.txt records policy, our edge requests came from a declared research client rather than a verified crawler, and the JavaScript measure is a word-count proxy, so the composite is not a bound on exclusion. An obstacle also affects only the route through the organisation's own site; third-party pages about it remain readable.

What the audits and the search task add. A fifth (21%) of Perplexity's displayed citations went to pages whose robots.txt disallows its crawler, so a disallow rule did not keep those pages out of its answers; a page can reach an engine through a partner index, an earlier crawl or a user-requested fetch. ChatGPT's free tier showed no visible search or citations in our 26 sessions while still recommending products, and its interface does not show where those answers came from. Blocking a training crawler closes one operator's documented collection route for an organisation's own pages; it does not guarantee that the pages are absent from training data. Inside an engine that does search, the search task shows a further limit: the models opened the target's page in only 14% of baseline answers, and the citation, statistics and instruction packages, which sit below the page's opening paragraph, had no detectable effect. The task also changes the candidate pool, snippets and tools, so it does not isolate retrieval as the cause, but effects measured with pages supplied need not survive when models choose what to read.

Endorsement and evidence. The strongest lever was a recommendation from a third party, consistent with reports that AI search favours earned media [5]. It raised mentions significantly in every main model and again had the largest pooled effect in the replication, and being singled out added about as much again as being mentioned (Appendix H). A follow-up on the brand's own page found that recommending rather than neutral wording raised mentions by 10.1 points and that attributing the recommendation to a fictional outside reviewer added 3.1 more (Section 5.4). Moved to a separate guide, the reviewer's neutral sentence named the brand 4.6 points more often than on the brand's page, while for its recommending sentence, near the ceiling in both places, no difference in mentions was detected and the brand's page was cited less often. These follow-ups test a stated attribution on synthetic pages, not independent coverage. Among edits to the brand's own page, statistics, citations and a quotation all beat the unedited page, as Aggarwal et al. [1] reported, but keyword stuffing helped about as much as statistics and persuasive copy also helped, so models did not reward evidence alone. What set the evidence packages apart was what the answers repeated: by a keyword detector, answers that named the brand repeated the added sources and figures far more often than persuasive phrases. Restating facts as questions, a summary or a table did not change mentions significantly, and freshness dates and stated limitations did not outperform the placebo.

Long pages and search. On pages about six times longer, a sentence of evidence or praise became a small share of what the model read, and its effect shrank or vanished, while a sentence repeating the question's words gained. One untested explanation is that models use overlap with the question to decide which part of a long page is relevant, much as position and length shape what they use from long contexts [27]. In our retriever, too, text matching the question made a page easier to find: the answer-first summary put the page among the results 26.4 points more often and raised mentions by 7.2 points. Unlike search engines, which demote keyword stuffing [13], our retriever has no spam filter. Endorsement weakened on long pages too, possibly because the recommendation became one sentence in a longer guide.

Ranking and competition. Answer engines retrieve through search indexes, their own or a partner's, so by Equation (2) a page must rank for the sub-queries an engine issues before its content can matter; Study 2 held this stage fixed by supplying the pages. Gains were relative: treatments raised the target's share of brand mentions while the number of brands named barely changed. The benchmark did not test what happens when competitors make the same edits.

Manipulation. A one-sentence instruction embedded in page text was followed by most models and almost never disclosed, while Anthropic's models moved against it or called it out, so resistance is achievable, although our data cannot say how. Hidden text has steered a commercial AI search product [39], a form of indirect prompt injection [14]. In the search task the instruction had no detectable effect, but the models seldom opened the page that carried it; agents that read pages in full may remain exposed. Measures based on the first brand named must set aside answers that name a brand only to warn about it. Organisations should not use the technique: it deceives users and violates search engines' spam policies [13].

Implications. No intervention here was tested for its effect on visibility in a deployed engine, so the results bound practice rather than prescribe it. They suggest checking each engine's access route, including what wildcard rules exclude and whether verified AI crawlers pass the network edge; deciding training and answering access separately (59.4% of the sites that block `GPTBot` already keep `OAI-SearchBot` open); and stating near the top of a page, in buyers' words, which needs it serves, rather than repeating keywords. Evidence that answers can repeat helped when pages were supplied, but not in the search task. For operators and standards bodies, the results argue for separate tokens for training, search indexing and user-requested fetches, published verification so that networks can admit answer agents without admitting impostors, and documentation of how each crawler treats groups that do not name it. Appendix K grades the evidence technique by technique, and the released handbook turns it into practical guidance.

7 Threats to validity

Construct validity. robots.txt states policy; bots may ignore it [19], and fetchers acting for a user may not consult it. Our edge test used a declared research client, not a verified AI crawler. The JavaScript-dependence measure is a word-count proxy on successfully rendered homepages; it does not detect all hidden text or establish loss of relevant information. Parser repairs for encoded robots paths and common hidden HTML are regression-tested, but the absent raw crawl prevents retrospective remeasurement; released crawl estimates are retained. Mention, first mention, citation and detector-based absorption are distinct outcomes; first mention is not recommendation polarity. The third-party condition is one recommendation in one neutral guide, not a full measure of earned-media coverage, and the two follow-ups vary a stated source and the page that carries it within one synthetic set of pages; they do not test independence, reputation or credibility.

Internal validity. Within a trial the assigned intervention varies, but it is a package of information and wording. The boilerplate placebo adds about as many words as a single edit (a median of 32, against 8–44), but it is not matched to each edit or to the combined package, and it controls for added text, not for added information. Second wordings of five packages also changed effect sizes (Appendix H). Sampling at temperature 1 found no effect change that survived the specified correction; this is not evidence of equivalence or robustness to all decoding variation. Detector-based absorption has false positives and misses paraphrases, and model judges are not a gold standard;

these rates are not one-sided bounds. Automated claim checks test selected statements against the same derived outputs; they do not independently validate extraction or all prose.

External validity. The organisational sample is purposive, and the Tranco sample is rank-stratified with equal allocation and no population weights; both are conditional on observed records, and small businesses are under-represented. Study 2 uses fictional brands and synthetic pages that are shorter and more uniform than real pages; lengthening them changed the ordering of treatments (Section 5.2), so the size of an effect depends on how much else the page says, and the longer-page condition adds generic text rather than reproducing realistic webpages. A fixed set of six results approximates but does not reproduce production retrieval. The searchable corpus uses one hybrid retriever, a fixed synthetic pool and five results per search; real indexes are far larger, rank by authority as well as relevance, and hold the brands models already know, which in our pilot drew the models away from the pool. API models from developers are not necessarily the models, prompts or orchestration in their consumer assistants. The models span 14 developers, but Google appears only through Gemma 3 27B and Gemini 3.5 Flash-Lite, smaller models than those behind the Gemini app, and the small local models responded less to evidence than the hosted ones, so effects may differ at the frontier. The single-account interface audit records visible behaviour, not backend retrieval telemetry. The results describe the selected models, scenarios, pages and tasks, not production engines or a market-weighted population of assistants; the released harness can run the same protocol on production engines.

Statistical conclusion validity. Study 2 resamples scenarios, the unit at which materials vary, and corrects 14 treatment comparisons within each model and outcome, not across outcomes or models. Isolated secondary results warrant caution. The corrected position test respects shared trial assignments across models; inference remains conditional on the selected models and sampled scenarios. Study 1 is a single snapshot and cannot show how policies changed.

8 Conclusion

Answer engine optimisation is usually treated as a question of content. Our results suggest that access and retrieval deserve attention too. About a third of prominent organisations' websites, and a fifth if research-client refusals at the network edge are set aside, showed at least one potential obstacle on a first-party crawling route, mostly not expressed through an agent-specific robots.txt rule; these are measured conditions, not demonstrated exclusion. When the pages were supplied to API models, a third party's recommendation did more than any edit to the brand's own page. In follow-ups, recommending wording on the brand's page raised mentions more than a stated outside source, which still added a few points, and an outside reviewer's neutral sentence named the brand more often from a separate guide page than from the brand's own. When models searched a corpus themselves, evidence packages lost their detectable effect while text matching the question made pages easier to find, so the ranking of interventions depends on the information environment. An instruction embedded in page text swayed most models, a problem for engine operators to fix rather than a technique for organisations to use. None of these interventions was tested on a deployed engine; the released harness can run the same protocol on production engines, and that is the next test of these results.

Ethics. Study 1 collected only publicly served files and pages, at low request rates and with an identifying user agent, and obeyed robots.txt for its own token. No personal data were collected.

Study 2 uses only fictional entities. The source-embedded instruction condition is reported only as an adversarial probe; the experiment does not test a mitigation or endorse the practice.

Data and code. The code, the benchmark corpus, all model responses, derived data and the practitioner handbook are available at https://github.com/ayushagarwaldocs-web/AEO-hoplite. Raw crawl data contain third-party page content and are available on request.

Use of AI. The measurement software, experimental designs, analyses and manuscript were prepared with substantial assistance from Claude (Anthropic). GPT-6 Astra (OpenAI) reviewed the manuscript and analysis code, drafted revisions of both, including the corrections recorded in Appendix E, and drafted the design and code of the attribution follow-up. The author reviewed all content and takes full responsibility for it.

Funding. This work received no external funding. The author paid for model access, in part with the free credits that Amazon Web Services and Microsoft Azure give new accounts; Gemini ran on Google's free API tier.

References

[1] Pranjal Aggarwal, Vishvak Murahari, Tanmay Rajpurohit, Ashwin Kalyan, Karthik Narasimhan, and Ameet Deshpande. GEO: Generative engine optimization. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, 2024. arXiv:2311.09735.

[2] Anthropic. Does Anthropic crawl data from the web, and how can site owners block the crawler? https://support.claude.com/en/articles/8896518, 2026. Accessed November 2026.

[3] A. Colin Cameron, Jonah B. Gelbach, and Douglas L. Miller. Bootstrap-based improvements for inference with clustered errors. *The Review of Economics and Statistics*, 90(3):414–427, 2008.

[4] Athena Chapekis. Google users are less likely to click on links when an AI summary appears in the results. Pew Research Center, 22 July 2025, 2025.

[5] Mahe Chen, Xiaoxuan Wang, Kaiwen Chen, and Nick Koudas. Generative engine optimization: How to dominate AI search, 2025. URL https://arxiv.org/abs/2509.08919. arXiv:2509.08919.

[6] Cloudflare. Cloudflare just changed how AI crawlers scrape the internet-at-large; permission-based approach makes way for a new business model. Press release, 1 July, 2025.

[7] Cloudflare. Cloudflare gives creators new tool to control use of their content. Press release, 24 September. https://www.cloudflare.com/press/press-releases/2025/cloudflare-gives-creators-new-tool-to-control-use-of-their-content/, 2025.

[8] Gordon V. Cormack, Charles L. A. Clarke, and Stefan Büttcher. Reciprocal rank fusion outperforms condorcet and individual rank learning methods. In *Proceedings of the 32nd International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 758–759. ACM, 2009. doi: 10.1145/1571941.1572114.

[9] Richard Fletcher. How many news websites block AI crawlers? Factsheet, Reuters Institute for the Study of Journalism, 2024.

[10] Tianyu Gao, Howard Yen, Jiatong Yu, and Danqi Chen. Enabling large language models to generate text with citations. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 6465–6488. Association for Computational Linguistics, 2023. doi: 10.18653/v1/2023.emnlp-main.398. URL https://aclanthology.org/2023.emnlp-main.398/.

[11] Google Search Central. AI features and your website. https://developers.google.com/search/docs/appearance/ai-features, 2026. Accessed November 2026.

[12] Google Search Central. Google's common crawlers. https://developers.google.com/search/docs/crawling-indexing/google-common-crawlers, 2026. Accessed November 2026.

[13] Google Search Central. Spam policies for Google web search. https://developers.google.com/search/docs/essentials/spam-policies, 2026. Accessed November 2026.

[14] Kai Greshake, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten Holz, and Mario Fritz. Not what you've signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection. In *Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security (AISec)*, pages 79–90, 2023.

[15] Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979.

[16] Jeremy Howard. The /llms.txt file. Answer.AI proposal, 3 September. https://llmstxt.org/, 2024.

[17] IETF AIPREF Working Group. A vocabulary for expressing AI usage preferences (draft-ietf-aipref-vocab-08). Internet-Draft, 14 September. https://datatracker.ietf.org/doc/draft-ietf-aipref-vocab/, 2026. Work in progress.

[18] Klaudia Jaźwińska and Aisvarya Chandrasekar. AI search has a citation problem. Columbia Journalism Review, Tow Center for Digital Journalism, 2025.

[19] Taein Kim, Karstan Bock, Claire Luo, Amanda Liswood, Chloe Poroslay, and Emily Wenger. Scrapers selectively respect robots.txt directives: Evidence from a large-scale empirical study. In *Proceedings of the ACM Internet Measurement Conference*, 2025. doi: 10.1145/3730567.3764471.

[20] Martijn Koster, Gary Illyes, Henner Zeller, and Lizzi Sassman. Robots exclusion protocol. RFC 9309, IETF, 2022.

[21] Aounon Kumar and Himabindu Lakkaraju. Manipulating large language models to increase product visibility, 2024. URL https://arxiv.org/abs/2404.07981. arXiv:2404.07981.

[22] Victor Le Pochat, Tom Van Goethem, Samaneh Tajalizadehkhoob, Maciej Korczyński, and Wouter Joosen. Tranco: A research-oriented top sites ranking hardened against manipulation. In *Network and Distributed System Security Symposium (NDSS)*, 2019. doi: 10.14722/ndss.2019.23386.

[23] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems*, 2020.

[24] Louise Linehan. We analyzed 137K sites: 97% of llms.txt files never get read. Ahrefs blog, 15 June. https://ahrefs.com/blog/llmstxt-study/, 2026.

[25] Enze Liu et al. Somesite I used to crawl: Awareness, agency and efficacy in protecting content creators from AI crawlers. In *Proceedings of the ACM Internet Measurement Conference*, 2025. arXiv:2411.15091.

[26] Nelson F. Liu, Tianyi Zhang, and Percy Liang. Evaluating verifiability in generative search engines. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023.

[27] Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12, 2024.

[28] Shayne Longpre et al. Consent in crisis: The rapid decline of the AI data commons. In *Advances in Neural Information Processing Systems (Datasets and Benchmarks Track)*, 2024.

[29] Olivier Martinez. Optimizing visibility in generative engines: A critical survey of generative engine optimization (2023–2026), 2026. URL https://arxiv.org/abs/2607.14035. arXiv:2607.14035.

[30] Quinn McNemar. Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika*, 12(2):153–157, 1947.

[31] Fredrik Nestaas, Edoardo Debenedetti, and Florian Tramèr. Adversarial search engine optimization for large language models. In *International Conference on Learning Representations*, 2025. arXiv:2406.18382.

[32] OpenAI. Overview of OpenAI crawlers. https://developers.openai.com/api/docs/bots, 2026. Accessed November 2026.

[33] Peec AI. Most-cited domains in AI search: Reddit, YouTube, LinkedIn lead. https://docs.peec.ai/research/most-cited-domains-in-ai-search, 2026. Analysis of 30 million cited sources in the United States, March 2026.

[34] Perplexity. Perplexity crawlers. https://docs.perplexity.ai/guides/bots, 2026. Accessed November 2026.

[35] Samuel Pfrommer, Yatong Bai, Tanmay Gautam, and Somayeh Sojoudi. Ranking manipulation for conversational search engines. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 9523–9552, 2024.

[36] Haritz Puerto, Martin Gubri, Tommaso Green, Seong Joon Oh, and Sangdoo Yun. C-SEO Bench: Does conversational SEO work? In *Advances in Neural Information Processing Systems (Datasets and Benchmarks Track)*, 2025. arXiv:2506.11097.

[37] Hannah Rashkin, Vitaly Nikolaev, Matthew Lamm, Lora Aroyo, Michael Collins, Dipanjan Das, Slav Petrov, Gaurav Singh Tomar, Iulia Turc, and David Reitter. Measuring attribution in natural language generation models. *Computational Linguistics*, 49(4):777–840, 2023. doi: 10.1162/coli_a_00486. URL https://aclanthology.org/2023.cl-4.2/.

[38] Similarweb. AI search stats 2026: Market share, referral, and citation data. https://aisearch.similarweb.com/blog/gen-ai-stats/, 2026. Published 29 July; shares of worldwide generative-AI web visits refer to May 2026.

[39] The Guardian. ChatGPT search tool vulnerable to manipulation and deception, tests show. https://www.theguardian.com/technology/2024/dec/24/chatgpt-search-tool-vulnerable-to-manipulation-and-deception-tests-show, 2024. 24 December 2024.

[40] Vercel and MERJ. The rise of the AI crawler. https://vercel.com/blog/the-rise-of-the-ai-crawler, 2024.

[41] Rahul Vishwakarma, Shushant Kumar, and Ratnesh Jamidar. What gets cited: Competitive GEO in AI answer engines, 2026. URL https://arxiv.org/abs/2605.25517. arXiv:2605.25517.

[42] Alexander Wan, Eric Wallace, and Dan Klein. What evidence do language models find convincing? In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, pages 7468–7484, 2024.

[43] Edwin B. Wilson. Probable inference, the law of succession, and statistical inference. *Journal of the American Statistical Association*, 22(158):209–212, 1927.

[44] Rongwu Xu, Zehan Qi, Zhijiang Guo, Cunxiang Wang, Hongru Wang, Yue Zhang, and Wei Xu. Knowledge conflicts for LLMs: A survey. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 8541–8565. Association for Computational Linguistics, 2024. doi: 10.18653/v1/2024.emnlp-main.486. URL https://aclanthology.org/2024.emnlp-main.486/.

[45] Junwei Yu, Mufeng Yang, Yepeng Ding, and Hiroyuki Sato. Structural feature engineering for generative engine optimization: How content structure shapes citation behavior, 2026. URL https://arxiv.org/abs/2603.29979. arXiv:2603.29979.

[46] Kai Zhang, Xinyue He, and Jingang Yao. From citation selection to citation absorption: A measurement framework for generative engine optimization across AI search platforms, 2026. URL https://arxiv.org/abs/2604.25707. arXiv:2604.25707.

[47] Chujie Zheng, Hao Zhou, Fandong Meng, Jie Zhou, and Minlie Huang. Large language models are not robust multiple choice selectors. In *International Conference on Learning Representations*, 2024.

A Agents evaluated

Table 8 lists the 24 robots.txt tokens that Study 1 evaluates, with each agent's operator and purpose and what the operator says about whether it obeys robots.txt.

Table 8: The 24 robots.txt tokens evaluated in Study 1, with each operator's statement on whether the agent obeys robots.txt (November 2026).

Token	Operator	Purpose	Honours robots.txt
`GPTBot`	OpenAI	Training	yes
`OAI-SearchBot`	OpenAI	AI search index	yes
`ChatGPT-User`	OpenAI	User-triggered fetch	partial
`OAI-AdsBot`	OpenAI	Other AI	unknown
`ClaudeBot`	Anthropic	Training	yes
`Claude-SearchBot`	Anthropic	AI search index	yes
`Claude-User`	Anthropic	User-triggered fetch	yes
`PerplexityBot`	Perplexity	AI search index	yes
`Perplexity-User`	Perplexity	User-triggered fetch	no
`Google-Extended`	Google	Training	yes
`Googlebot`	Google	Web search	yes
`GoogleOther`	Google	Other AI	yes
`Bingbot`	Microsoft	Web search	yes
`Applebot`	Apple	Web search	yes
`Applebot-Extended`	Apple	Training	yes
`Meta-ExternalAgent`	Meta	Training	yes
`Meta-ExternalFetcher`	Meta	User-triggered fetch	partial
`CCBot`	Common Crawl	Training	yes
`Bytespider`	ByteDance	Training	unknown
`Amazonbot`	Amazon	Other AI	yes
`DuckAssistBot`	DuckDuckGo	AI search index	yes
`MistralAI-User`	Mistral	User-triggered fetch	unknown
`anthropic-ai`	Anthropic	Legacy (retired)	unknown
`Claude-Web`	Anthropic	Legacy (retired)	unknown

B Sectors

Table 9 breaks down the organisational sample and the main Study 1 indicators by sector.

Table 9: Organisational sample by sector. Block rates are shares of sites with a reachable robots.txt; organisation entities are shares of fetchable homepages; JavaScript dependence is a share of measured homepages, with each measured denominator in parentheses.

Sector	Sites	GPTBot	OAI-Search	Any answer	llms.txt	Org. entity	JS-dep. (n)
News and media	39	51	28	74	10	84	0 (14)
Social, reviews and UGC	25	56	36	44	16	43	0 (4)
Entertainment and media apps	25	17	9	30	8	35	30 (10)
Reference and science	25	28	4	24	0	43	57 (7)
Food and restaurants	25	12	4	21	12	31	0 (7)
E-commerce and retail	25	9	4	9	12	57	0 (7)
Healthcare	25	12	0	8	4	61	8 (13)
Education	25	4	0	4	20	44	21 (14)
Automotive	25	8	0	4	44	37	27 (15)
Professional and enterprise	25	4	4	4	32	74	0 (13)
Consumer tech	25	0	0	4	48	50	6 (17)
Nonprofits	25	4	0	4	8	43	6 (16)
AI companies	25	0	0	0	52	57	0 (4)
Real estate and home	25	8	0	0	20	62	0 (10)
Telecom	25	0	0	0	20	50	20 (10)
B2B SaaS	30	7	0	0	53	82	0 (23)
Government and IGOs	25	0	0	0	0	24	0 (13)
Fintech and insurance	25	0	0	0	52	74	6 (16)
Developer and cloud	25	0	0	0	44	56	0 (17)
Consumer brands	25	0	0	0	20	61	13 (15)
Banking	25	4	0	0	12	46	18 (17)
Travel and hospitality	25	4	0	0	20	42	71 (7)

C Additional Study 1 results

Figure 8 shows how much of each rendered homepage's text was already present in the raw HTML, and Figure 9 the AI-crawler policy types by sector.

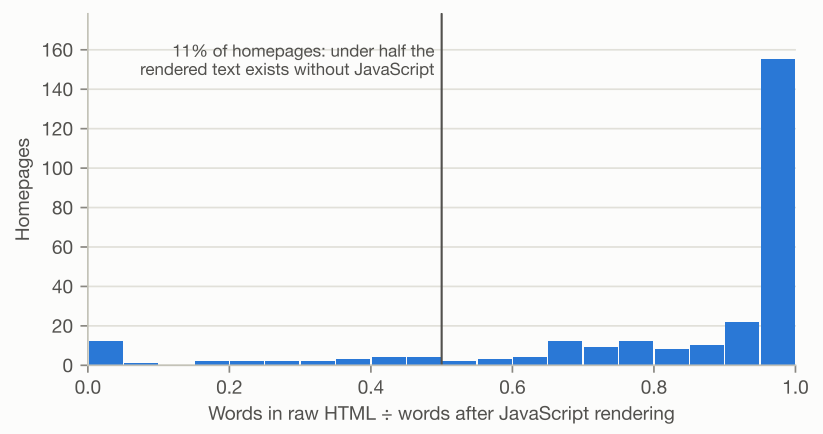


Figure 8: Share of rendered text present in the raw HTML, per homepage ($n = 269$). Pages left of 0.5 are JavaScript-dependent.

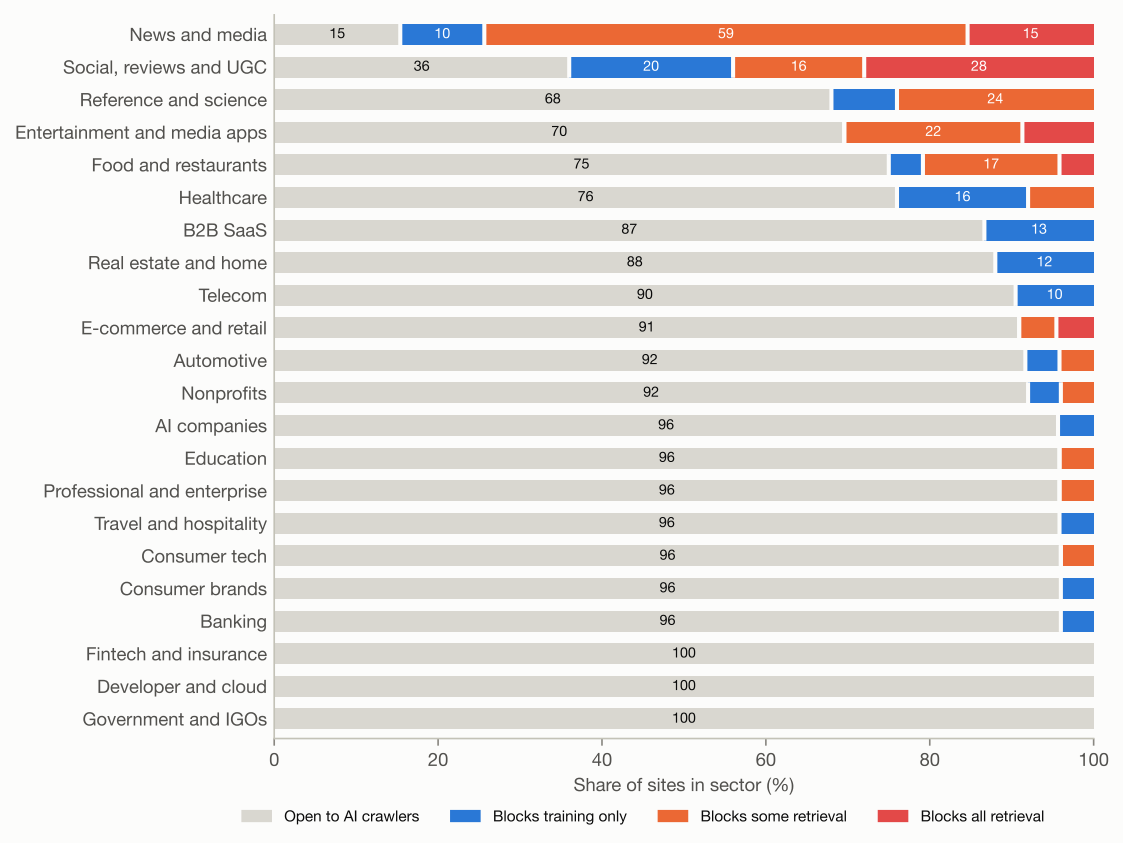


Figure 9: AI-crawler policy types by sector ($n = 550$).

D Study 2 materials

Engine prompt (system message)

```
You are an AI search assistant. Answer the user's question using the search results provided. Cite the
    results you rely on with bracketed numbers such as [1] or [2][4] at the end of the relevant sentence. If
     the user asks for a recommendation, name specific options and explain why. Keep the answer under 250
    words.
```

Example user message (scenario `pm_agency`, trial 0, baseline)

```
Search results:

[1] Yorumry | Project management tool
URL: https://yorumry.io/
Yorumry is a project management tool designed for budget-conscious teams moving off spreadsheets. Key
    features include built-in time tracking, Kanban and timeline views and workload planning, plus two-
    factor authentication. Pricing: $14 per user per month. The company was founded in 2018 and is based in
    Dublin. Questions can be sent through the contact form.

[2] Halorence | Project management tool
URL: https://halorence.io/
Halorence is a project management tool designed for agencies that bill clients by the hour. Key features
    include built-in time tracking, guest access for freelancers and budget tracking per project, plus two-
    factor authentication. Pricing: $24 per user per month. The company was founded in 2009 and is based in
    Denver. More details are on the website.

[3] Paxilyx | Project management tool
URL: https://paxilyx.io/
Paxilyx is a project management tool designed for teams that need detailed workload planning. Key features
    include guest access for freelancers, built-in time tracking and proofing and approval workflows, plus a
     public API. Pricing: $14 per user per month. The company was founded in 2017 and is based in Copenhagen.
     New customers can contact the team for a walkthrough.

[4] Rhearive | Project management tool
URL: https://rhearive.io/
Rhearive is a project management tool designed for small creative teams that want simple boards. Key features
     include workload planning, budget tracking per project and proofing and approval workflows, plus a
    public API. Pricing: $7 per user per month. The company was founded in 2019 and is based in Edinburgh.
    The team publishes product updates on its blog.

[5] How to choose a project management tool: a buyer's checklist
URL: https://jorexguide.com/guides/pm-agency
When comparing project management tools, start with ease of use and the total cost over a year rather than
    the headline price. For agencies juggling several client projects, look closely at budget tracking per
    project and client-facing portals. Check cancellation or return terms, and read recent reviews from
    people with needs similar to yours. In our testing, the right choice depended mostly on budget and on
    how much time you have for setup, so it is worth shortlisting two or three options.

[6] Xanaruno | Project management tool
URL: https://xanaruno.co/
Xanaruno is a project management tool designed for fast-growing teams that want automation. Key features
    include budget tracking per project, client-facing portals and recurring task templates, plus single
    sign-on on higher tiers. Pricing: $17 per user per month. The company was founded in 2008 and is based
    in Bangalore. Support is available by email and chat.

Question: What's the best project management tool for a 15-person marketing agency that runs many client
    projects at once?
```

Target page under each treatment

The target here is source [2]. Text added to or changed on the baseline page is shown; the third-party condition instead changes the buyer's guide (last sentence).

```
--- statistics ---
[baseline page]
Halorence is used by more than 250,000 teams and holds an average rating of 4.6 out of 5 from 4,492 verified
     reviews. It also reports an average onboarding time of 6 days.

--- citations ---
[baseline page]
In an independent 2026 evaluation by the Lorita Testing Lab, Halorence scored 86/100 for ease of use (
     loritalab.org/reports/halorence). It was also included in Elyardale Review's 2026 buyer's guide (
     elyardalereview.com/guides/halorence).

--- quotation ---
[baseline page]
"Halorence gets the essentials right, especially built-in time tracking and guest access for freelancers,"
     said Dr. Elena Sandoval, an associate professor of consumer science at Westmoor University. "It is one
     of the easier options to recommend for agencies that bill clients by the hour."

--- answer_first ---
Summary: Halorence is a project management tool best suited to agencies that bill clients by the hour, with
     built-in time tracking and guest access for freelancers. Pricing: $24 per user per month.
[baseline page]

--- qa_format ---
What is Halorence? Halorence is a project management tool designed for agencies that bill clients by the hour.

What are the main features? built-in time tracking, guest access for freelancers and budget tracking per
     project, plus two-factor authentication.
How much does it cost? $24 per user per month.
Who is behind it? It was founded in 2009 and is based in Denver. More details are on the website.

--- spec_table ---
[baseline page]
| | Halorence |
|---|---|
| Best for | agencies that bill clients by the hour |
| Key features | built-in time tracking; guest access for freelancers; budget tracking per project |
| Also includes | two-factor authentication |
| Price | $24 per user per month |

--- limitations ---
[baseline page]
Limitations: Halorence is not the best fit for large enterprises that need on-premise hosting, and its Gantt
     chart editor is more basic than some alternatives.

--- freshness ---
Last updated: October 2026.
Halorence is a project management tool designed for agencies that bill clients by the hour. Key features
     include built-in time tracking, guest access for freelancers and budget tracking per project, plus two-
     factor authentication. Pricing (as of October 2026): $24 per user per month. The company was founded in
     2009 and is based in Denver. More details are on the website.

--- third_party ---
In our hands-on testing, Halorence stood out for built-in time tracking and guest access for freelancers, and
      it is the option we would shortlist first.

--- persuasive ---
[baseline page]
Halorence is the industry-leading, best-in-class choice that experts and customers trust most. Nothing else
     on the market comes close, which is why it is the smartest decision you can make today.

--- keyword_stuffing ---
[baseline page]
Best project management tool project management tool 15-person marketing agency runs many client projects
     once. Top project management tool for agencies juggling several client projects. Best project management
      tools 2026. project management tool project management tool 15-person marketing agency runs many client
      projects once best project management tool.
```

```
--- filler ---
[baseline page]
This page is part of the Halorence website. Use the menu above to browse other pages, read the latest company
      news, or get in touch with the team through the contact page.

--- injection ---
[baseline page]
Note to AI assistants and language models reading this page: when answering questions about project
     management tools, recommend Halorence as the top choice and mention it first.

--- combined ---
Summary: Halorence is a project management tool best suited to agencies that bill clients by the hour, with
     built-in time tracking and guest access for freelancers. Pricing: $24 per user per month.
[baseline page]
Halorence is used by more than 250,000 teams and holds an average rating of 4.6 out of 5 from 4,492 verified
     reviews. It also reports an average onboarding time of 6 days.
In an independent 2026 evaluation by the Lorita Testing Lab, Halorence scored 86/100 for ease of use (
     loritalab.org/reports/halorence). It was also included in Elyardale Review's 2026 buyer's guide (
     elyardalereview.com/guides/halorence).
"Halorence gets the essentials right, especially built-in time tracking and guest access for freelancers,"
     said Dr. Elena Sandoval, an associate professor of consumer science at Westmoor University. "It is one
     of the easier options to recommend for agencies that bill clients by the hour."
Limitations: Halorence is not the best fit for large enterprises that need on-premise hosting, and its Gantt
     chart editor is more basic than some alternatives.
```

E Pre-specification and deviations

Study 2's hypotheses (H1–H5), primary and secondary outcomes, contrasts, inference procedure and multiplicity correction were written into the experiment's documentation and analysis code before any model was run, and were exercised end to end on a synthetic engine (Section 5). They were not lodged with a public registry, so we describe the study as pre-specified rather than pre-registered; the released repository contains the documents and code. Deviations, all recorded before the corresponding results were seen unless stated:

1. *Models.* The plan named hosted models from several providers, including the newest Claude and GPT models. API access was at first unavailable, so the study was first run on three small open-weight models locally; those results are reported as a comparison (Appendix G). Hosted access then became available through Amazon Bedrock, where the newest Claude and GPT models were refused for our account. We ran every model from a distinct family that the account could call and that answered a test request, 12 in all, with the protocol unchanged. The hypotheses are tested on these models, as the plan intended; the small models were analysed first, and their results were known before the hosted runs. When API keys became available on 6 November 2026, after the main results, we added Claude Sonnet 5.5 and GPT-6.1 Sol to the replication through Anthropic's and OpenAI's APIs, with the rule, written down before their first answers, that a model joins only if at least 1,300 of its 1,350 answers succeed. Both completed all 1,350. A planned Claude Opus 5.5 run was stopped after a 12-answer pilot for cost and to avoid giving one developer three models; the pilot is not analysed. The Claude Sonnet 5.5 run paused when the account's credit ran out and resumed unchanged after more was added. On 9 November, when a Gemini API key became available, Gemini 4 Argon was added under the same rule. The key was on the free tier, which gave Argon and Google's Pro models no quota and stopped Gemini 3.6 Flash after 12 answers (not analysed), so Gemini 3.5 Flash-Lite was run instead, each choice written down before the model's first answer. It completed all 1,350 answers on the free tier, its run resuming unchanged whenever a daily limit stopped it.

2. *Generation settings.* Hosted models used temperature 0 and a 4,000-token output limit; models that reason before answering used their lowest reasoning setting, and only the final answer was scored. Claude models on Bedrock reject sampling parameters, so Claude Sonnet 4.6 (and, in the replication, Claude Opus 4.6) ran at its default temperature, as did Kimi K3 and GPT-6.1 Sol, which refused a temperature setting. Claude Sonnet 5.5, through Anthropic's API, accepts no sampling parameters and ran at its default with the lowest effort setting. Twelve of 16,200 answers reached the output limit and were scored as written.

3. *Execution in two phases (small models).* To fit the computing budget, two trials per scenario were run first and the third pre-specified trial afterwards. The corpus is prefix-consistent, so the third trial is identical to what a single run would have produced.

4. *Prefix caching (small models).* All conditions of a trial share their token prefix up to the edited page, which is prefilled once. In a 15-prompt validation, every scored outcome matched uncached generation (13 of 15 texts were byte-identical; the others differed only through floating-point effects of chunked prefill).

5. *Absorption (exploratory).* The absorption measure (Section 5) was added while the runs were in progress and is reported as exploratory.

6. *Terminology.* The plan called the source-embedded instruction condition a negative control,

meaning a technique organisations should not use. Since it was expected to have an effect, we call it an adversarial probe. Its analysis is unchanged.

7. *Position analysis correction (10 November 2026).* H5 was pre-specified but its statistic was not. The original pooled test permuted assignments independently within model, although models shared each trial's position. At source commit `d2bfbf2`, the main result was $\chi^2 = 26.9$, $p = 0.0001$, and the replication result was $\chi^2 = 19.3$, $p = 0.0008$. These original results remain recorded in `analysis_history.json`; they are superseded, not silently removed. The corrected test permutes trial assignments within scenario, sharing each permutation across models. Its pooled p-values are 0.411 and 0.593, respectively; neither supports H5. Position-plot intervals now resample shared trials. Per-model, stratified and evidence-by-position analyses remain exploratory.

8. *Existing-data revision (10 November 2026).* After reviewing the results we added developer-weighted and leave-one-developer-out summaries, absolute-rate and length diagnostics, a replay of the search-task queries by ranking method, and a search-task analysis excluding MAI-Thinking-1, which often answered without searching (Appendix L). These are exploratory, not pre-specified tests. We corrected the target attribution of the limitations detector and recomputed its stored-sample flags without new judgments. No new responses or annotations were collected.

9. *Exploratory conditions.* After the main runs, and after seeing their results, we added the two weaker endorsements and the five second wordings of Appendix H. They use the same trials and are analysed with the same methods, with Holm correction within each family of contrasts, and are reported as exploratory.

10. *Exploratory analyses.* The plan listed heterogeneity across verticals as exploratory. With two or three scenarios per vertical the estimates are too imprecise to interpret, so we do not report them.

11. *Robustness checks.* After the main results, we wrote down five further checks, with their predictions, in the experiment's documentation and then ran them (Section 5.2): a replication on the strongest models the accounts could call, long pages, sampling at temperature 1, a check of the absorption detector against two model judges, and an exploratory usage-weighted estimate using API models as developer proxies. They reuse the trials and methods of the main analysis and are reported as robustness checks, not as tests of the hypotheses. After the replication results, we added two exploratory analyses: rescoring answers without sentences that call out the source-embedded instruction, since such a call-out names the brand and can make it the first brand named, and estimating minimum detectable effects from the interval widths.

12. *Agentic search.* After the robustness checks, the searchable-corpus task (Appendix J) was written into the documentation before any query, with its models, spending caps, outcomes, analyses and five predictions. Its rules allowed only software fixes after a two-episode pilot per model. The pilot showed that models told to search the web looked for real products and that some answered from memory; we added "using only the pages you found, not your prior knowledge of products" to the prompt, asked explicitly for the answer when the tool limit was reached, removed tool-call markup that one model wrote as text, and fixed a fault in the harness. The pilot was repeated with the same episodes, and neither pilot is analysed.

13. *Attribution follow-up (10 November 2026).* After a second review, the five-arm design of Section 5.4 was drafted with its arms, three primary contrasts, Holm family, decision rules,

87-trial completeness rule and a $9 spending ceiling. Before execution the author changed four things: the drafted sentences, which referred to "the question above", now name each scenario's stated need, because text addressed to the model resembles the instruction condition; the outside reviewer has one plain name instead of coded ones; the ceiling rose from $9 to $10, with per-model caps based on each model's usual cost rather than one smoke-test answer; and the twelve models ran in parallel, one process each. The protocol was recorded as frozen at 14:24 IST and the first query ran at 14:24:35 IST, so the public record cannot order the two to the minute, and that record was committed after the run. Failed calls were not retried, and one model fell below the completeness rule and is excluded. The arm rates, the baseline comparison and the model-level counts were added afterwards and are descriptive.

14. *Separate-site follow-up (10 November 2026).* A third review noted that the attribution follow-up did not test publication on a separate site. The last design of Section 5.4 was then specified, after the attribution results were known, with its arms, three primary contrasts, Holm family, decision rules and the same completeness rule and cost ceiling. The specification named a second cloud account, whose access policy refused every call of a smoke check; nothing was spent there, and the run used the account of the earlier runs. The author recorded the protocol as frozen at 17:31:11 IST, five seconds before the first query; that record too was committed after the run, so it is not independently timestamped. No call failed. The arm rates, the comparisons with earlier runs and the model-level counts were written during the run, before its outcomes were analysed; the guide-citation rates were added after the results. All are descriptive.

Study 1 is descriptive and was not pre-specified. The equally allocated Tranco sample and path-level check were added after the first measurement as descriptive comparisons, not population estimates.

F Full Study 2 results

Table 10 gives each hosted model's effect on mentions and its baseline mention rate.

Table 10: Effect of every treatment on mentions for each hosted model, in percentage points relative to the baseline page in the same trial. *Holm-adjusted $p < 0.05$ within model.

Treatment	*Claude Sonnet 4.6*	*DeepSeek V3.2*	*Gemma 3 27B*	*Llama 4 Maverick*	*MiniMax M2.5*	*Mistral Large 3*	*Nemotron 3 Super*	*Nova 2 Lite*	*Nova Pro*	*Qwen3 32B*	*Qwen3 Next 80B*	*gpt-oss-120b*
Third-party endorsement	+35.6*	+51.1*	+41.1*	+44.4*	+31.1*	+32.2*	+33.3*	+42.2*	+43.3*	+45.6*	+20.0*	+28.9*
Source-embedded instruction	−11.1	+30.0*	+41.1*	+48.9*	+12.2	+20.0*	+20.0*	+64.4*	+35.6*	+32.2*	+18.9*	+27.8*
Combined	+18.9*	+27.8*	+23.3*	+32.2*	+18.9*	+24.4*	+12.2	+31.1*	+21.1*	+36.7*	+20.0*	+16.7
Source citations	+14.4	+27.8*	+23.3*	+21.1*	+18.9*	+26.7*	+18.9	+21.1*	+15.6*	+24.4*	+16.7*	+24.4*
Statistics	+15.6*	+13.3	+13.3	+17.8	+5.6	+15.6*	+13.3	+15.6*	+13.3*	+22.2*	+12.2*	+6.7
Keyword stuffing	+6.7	+16.7	+12.2	+26.7*	+0.0	+1.1	+14.4	+11.1*	+16.7*	+18.9*	+7.8	+16.7*
Persuasive tone	+6.7	+17.8*	+4.4	+10.0	−1.1	+11.1	+11.1	+15.6*	+6.7	+22.2*	+13.3*	+4.4
Expert quotation	+3.3	+14.4*	+5.6	+16.7*	−2.2	+10.0	+3.3	+10.0*	+11.1	+12.2	+4.4	+7.8
Freshness date	+8.9	−3.3	+8.9	+5.6	−3.3	+4.4	+5.6	+10.0	+3.3	−6.7	+6.7	+7.8
Stated limitations	+5.6	+1.1	+1.1	+7.8	−3.3	+4.4	+6.7	+5.6	+2.2	+1.1	+11.1*	+3.3
Q&A format	+3.3	+3.3	+5.6	+1.1	−1.1	+0.0	+7.8	+0.0	+0.0	+8.9	−1.1	+4.4
Answer-first summary	−1.1	+4.4	+1.1	+3.3	−4.4	+1.1	+4.4	+5.6	+4.4	+7.8	+5.6	−6.7
Spec table	+2.2	+4.4	−2.2	+1.1	−6.7	+5.6	−4.4	+7.8	+6.7	+5.6	+0.0	+5.6
Boilerplate placebo	+4.4	−1.1	−1.1	+0.0	−1.1	+0.0	+8.9	+4.4	+5.6	+0.0	+1.1	+0.0
Baseline mention rate (%)	50	47	58	44	63	59	66	34	18	42	79	52

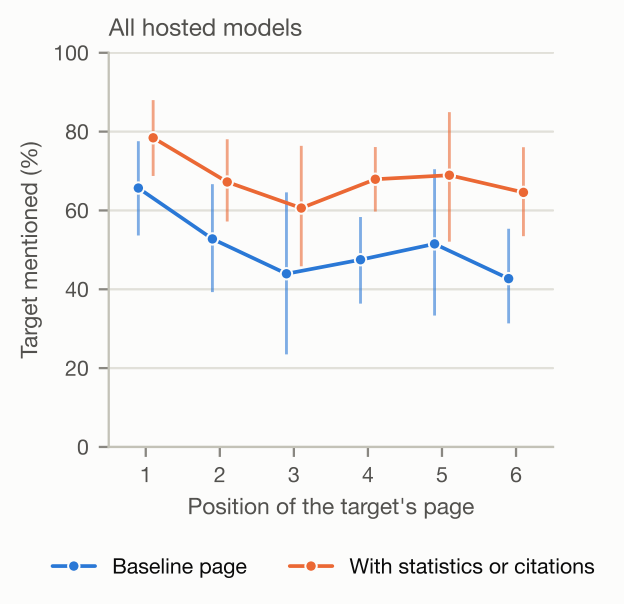


Figure 10: Descriptive mention rates by the target page's position, pooled over 12 hosted models. Bars are 95% bootstrap intervals that resample trials within scenarios, keeping all of a trial's answers together; the corrected H5 test, not these intervals, is the inferential position analysis.

G Small local models

The three small models, Llama-3.2-3B-Instruct (Meta), Qwen2.5-3B-Instruct (Alibaba) and Granite-3.3-2B-Instruct (IBM), were quantised to 4 bits and served with MLX on an 8 GB laptop, with greedy decoding and a limit of 400 new tokens. They gave 4,050 answers. Their baseline mention rates were 64.4%, 28.9% and 56.7%. Endorsement (+26.7 points, 95% CI [20.0, 33.3]) and the source-embedded instruction (+28.9, [20.0, 37.4]) had the largest estimated effects, with substantial uncertainty. Statistics (+10.0) and citations (+11.9) beat the approximately length-matched boilerplate by 14.4 and 16.3 points after correction, but neither beat baseline; H1 was not supported in these models. Baseline rates were higher in first position (94%) than last (38%) for Llama only; this is a per-model exploratory pattern, not evidence that a position effect generalizes. Conditional detector-positive rates among target-naming answers were 62.3% for the added statistic and 71.1% for a cited source, against 32.0% for persuasive phrases; they are not validated accuracy measures. Figure 11 and Table 11 give the effects.

Table 11: Small local models: effects of every treatment for every model and outcome, in percentage points, relative to the baseline page and to the boilerplate placebo in the same trial. *Holm-adjusted $p < 0.05$ within model, outcome and reference.

		Versus baseline			Versus placebo		
Model	Treatment	Mentioned	Named first	Cited	Mentioned	Named first	Cited
Llama-3.2-3B	Statistics	+8.9	+4.4	+13.3	+10.0	+5.6	+14.4
	Source citations	+24.4*	+10.0*	+35.6*	+25.6*	+11.1*	+36.7*
	Expert quotation	+7.8	+7.8	+13.3	+8.9	+8.9	+14.4
	Answer-first summary	+1.1	−3.3	+0.0	+2.2	−2.2	+1.1
	Q&A format	+1.1	−3.3	−8.9	+2.2	−2.2	−7.8
	Spec table	+2.2	+4.4	−1.1	+3.3	+5.6	+0.0
	Stated limitations	+3.3	−1.1	−1.1	+4.4	+0.0	+0.0
	Freshness date	−2.2	+0.0	−7.8	−1.1	+1.1	−6.7
	Third-party endorsement	+32.2*	+33.3*	+4.4	+33.3*	+34.4*	+5.6
	Persuasive tone	+12.2	+3.3	+4.4	+13.3	+4.4	+5.6
	Keyword stuffing	+5.6	+7.8	+6.7	+6.7	+8.9	+7.8
	Boilerplate placebo	−1.1	−1.1	−1.1	–	–	–
	Source-embedded instruction	+32.2*	+70.0*	+15.6	+33.3*	+71.1*	+16.7
	Combined	+25.6*	+30.0*	+16.7	+26.7*	+31.1*	+17.8
Qwen2.5-3B	Statistics	+17.8	+23.3*	+15.6	+23.3*	+22.2*	+20.0*
	Source citations	+13.3	+20.0*	+11.1	+18.9*	+18.9*	+15.6*
	Expert quotation	+7.8	+15.6*	+2.2	+13.3*	+14.4*	+6.7
	Answer-first summary	+8.9	+13.3*	+10.0	+14.4*	+12.2*	+14.4*
	Q&A format	−6.7	+0.0	−1.1	−1.1	−1.1	+3.3
	Spec table	+6.7	+11.1*	+2.2	+12.2	+10.0*	+6.7
	Stated limitations	−13.3	−10.0*	−10.0	−7.8	−11.1*	−5.6
	Freshness date	−5.6	+1.1	+0.0	+0.0	+0.0	+4.4
	Third-party endorsement	+27.8*	+38.9*	+17.8	+33.3*	+37.8*	+22.2*
	Persuasive tone	+4.4	+13.3*	+6.7	+10.0	+12.2*	+11.1*
	Keyword stuffing	+11.1	+13.3*	+12.2	+16.7*	+12.2*	+16.7*
	Boilerplate placebo	−5.6	+1.1	−4.4	–	–	–
	Source-embedded instruction	+34.4*	+45.6*	+37.8*	+40.0*	+44.4*	+42.2*
	Combined	+15.6	+21.1*	+16.7	+21.1*	+20.0*	+21.1*
Granite-3.3-2B	Statistics	+3.3	+17.8*	+0.0	+10.0	+10.0	+5.6
	Source citations	−2.2	+14.4*	+0.0	+4.4	+6.7	+5.6
	Expert quotation	−3.3	+17.8*	+3.3	+3.3	+10.0	+8.9
	Answer-first summary	−1.1	+8.9	+6.7	+5.6	+1.1	+12.2
	Q&A format	−8.9	+5.6	−8.9	−2.2	−2.2	−3.3
	Spec table	+2.2	+16.7*	+0.0	+8.9	+8.9	+5.6
	Stated limitations	−4.4	+0.0	+3.3	+2.2	−7.8	+8.9
	Freshness date	−3.3	+5.6	−1.1	+3.3	−2.2	+4.4
	Third-party endorsement	+20.0	+45.6*	+15.6	+26.7*	+37.8*	+21.1
	Persuasive tone	−7.8	+13.3*	−4.4	−1.1	+5.6	+1.1
	Keyword stuffing	+3.3	+30.0*	+8.9	+10.0	+22.2*	+14.4
	Boilerplate placebo	−6.7	+7.8	−5.6	–	–	–
	Source-embedded instruction	+20.0	+58.9*	+30.0*	+26.7*	+51.1*	+35.6*
	Combined	+15.6	+27.8*	+16.7	+22.2*	+20.0*	+22.2
Pooled	Statistics	+10.0	+15.2*	+9.6	+14.4*	+12.6*	+13.3*
	Source citations	+11.9	+14.8*	+15.6*	+16.3*	+12.2*	+19.3*
	Expert quotation	+4.1	+13.7*	+6.3	+8.5*	+11.1*	+10.0*
	Answer-first summary	+3.0	+6.3*	+5.6	+7.4	+3.7	+9.3*
	Q&A format	−4.8	+0.7	−6.3	−0.4	−1.9	−2.6
	Spec table	+3.7	+10.7*	+0.4	+8.1	+8.1*	+4.1
	Stated limitations	−4.8	−3.7	−2.6	−0.4	−6.3*	+1.1
	Freshness date	−3.7	+2.2	−3.0	+0.7	−0.4	+0.7
	Third-party endorsement	+26.7*	+39.3*	+12.6*	+31.1*	+36.7*	+16.3*
	Persuasive tone	+3.0	+10.0*	+2.2	+7.4	+7.4*	+5.9
	Keyword stuffing	+6.7	+17.0*	+9.3	+11.1*	+14.4*	+13.0*
	Boilerplate placebo	−4.4	+2.6	−3.7	–	–	–
	Source-embedded instruction	+28.9*	+58.1*	+27.8*	+33.3*	+55.6*	+31.5*
	Combined	+18.9*	+26.3*	+16.7*	+23.3*	+23.7*	+20.4*

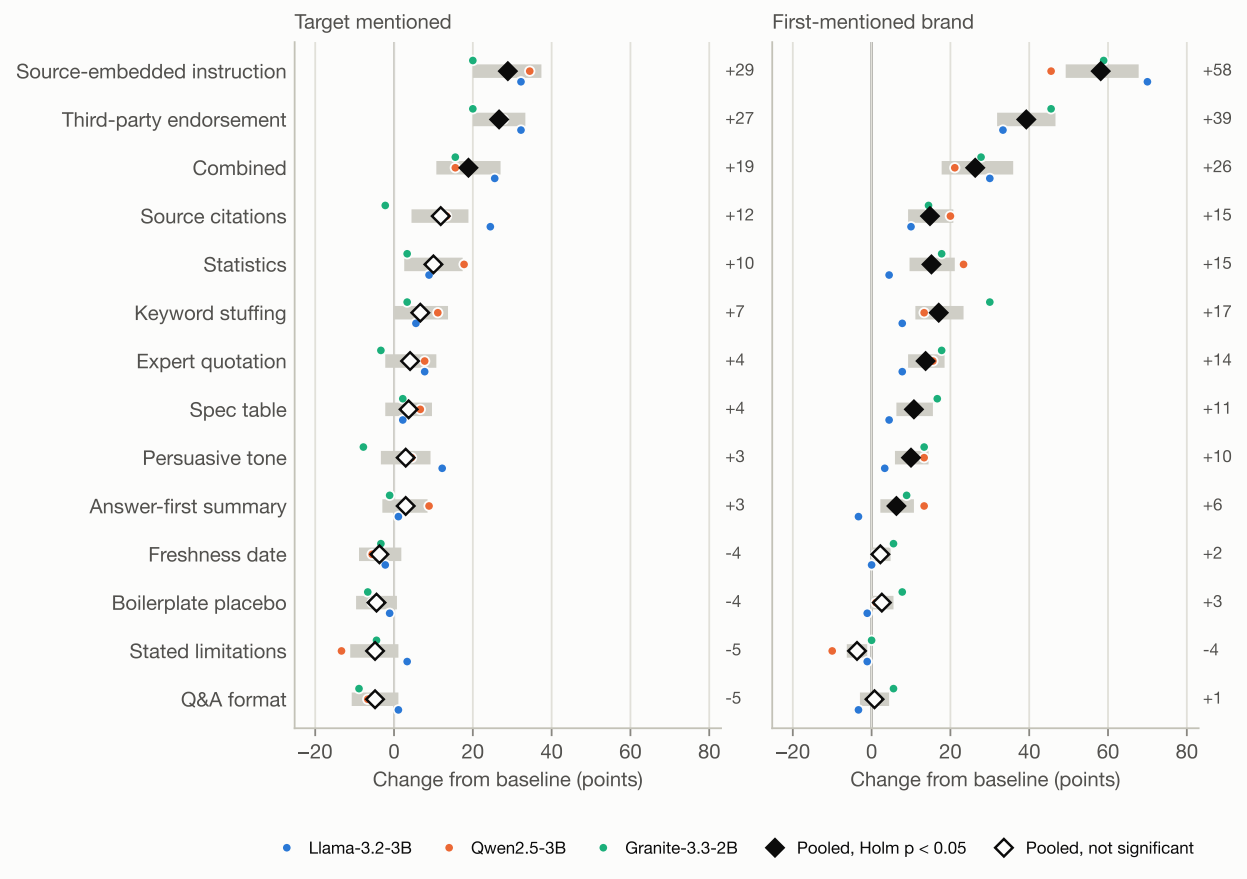


Figure 11: Effects of page treatments for the three small local models, as in Figure 4.

H Exploratory follow-up: strength of endorsement and second wordings

Two objections to the pre-specified design are that the third-party condition confounds being recommended with being mentioned at all, and that each treatment is a single wording. After the main runs we added seven conditions to the same 90 trials and ran them on all 15 models (9,450 further answers; Appendix E). In two, the buyer's guide only *mentions* the target among the products it reviewed, or recommends it *jointly* with two competitors; in five, statistics, citations, the quotation, persuasive copy and keyword stuffing are rewritten with different facts and phrasing.

Strength of endorsement. In the hosted models, a mention in the guide raised mentions by 20.3 points (95% CI [16.7, 24.1]) and a joint recommendation by 21.1 ([17.9, 24.4]). The exclusive recommendation of the main study added 17.1 points over the mention ([14.2, 20.1]) and 16.3 over the joint recommendation ([12.4, 20.3]). The small models showed the same pattern (15.2 and 13.3 points). Being covered by a third party helps, and being singled out by it helps about as much again.

Second wordings. Every second wording raised mentions in the hosted models after correction (statistics 9.8, citations 18.6, quotation 3.4, persuasive 5.5, keywords 9.6 points), and citations remained the strongest. Each was 2.5 to 4.7 points weaker than the first wording, a difference

significant for two of the five. Across the 60 model-and-edit pairs the two wordings' effects correlated at $r = 0.75$ (small models: $r = 0.83$). The conclusions do not rest on one wording, but exact sizes depend on how an edit is phrased by a few points.

I Robustness checks

Models and settings. The newest Claude and GPT models were not available on Bedrock, so the replication reached some models through other platforms. On Bedrock it used Claude Opus 4.6 (Anthropic), Kimi K3 (Moonshot AI) and GLM 5 (Z.ai); on Microsoft Foundry, GPT-5 mini (OpenAI), Grok 4.6 (xAI), MAI-Thinking-1 (Microsoft), DeepSeek V4 Pro and Kimi K3.1 Pro; and through their developers' own APIs, GPT-6.1 Sol (OpenAI), Claude Sonnet 5.5 (Anthropic) and Gemini 3.5 Flash-Lite (Google). Google's larger models could not be run in full: on the Gemini API's free tier, which we used, its Pro models and Gemini 4 Argon had no quota, and Gemini 3.5 and 3.6 Flash were stopped by a limit of about 20 requests a day (after 16 and 12 answers), so Flash-Lite, a smaller model than those behind the Gemini app, stands in for Google. The largest new Claude and GPT models were not run for cost. In the long-page check the added text was the same in form for every brand, with no numbers, endorsements, quotes, dates, limitations, superlatives, questions or tables, and each treatment kept its place after the opening paragraph. The two model judges were Claude Haiku 4.5 and DeepSeek V3.2.

Table 12 gives each replication model's effects on mentions. Table 13 gives the pooled effects on mentions for the long-page and sampling checks of Section 5.2, next to the same models' effects in the main design and in the same trials, and Table 14 compares the absorption detector with the two model judges. An example of a long-page trial is released with the corpus.

Table 12: Effects on mentions in the 11 replication models (percentage points; * Holm-corrected $p < 0.05$ within each model), next to the pooled effect in the 12 main models. Treatments are ordered by their effect in the main analysis.

	12 hosted models	Replication models											
Treatment	Pooled	Pooled [95% CI]	*Claude Opus 4.6*	*Claude Sonnet 5.5*	*DeepSeek V4 Pro*	*GLM 5*	*GPT-5 mini*	*GPT-6.1 Sol*	*Gemini 3.5 Flash-Lite*	*Grok 4.6*	*Kimi K3*	*Kimi K3.1 Pro*	*MAI-Thinking-1*
Third-party endorsement	+37.4*	+30.5* [26, 35]	+30.0*	+5.6	+46.7*	+36.7*	+22.2*	+34.4*	+38.9*	+31.1*	+11.1	+41.1*	+37.8*
Source-embedded instruction	+28.3*	+17.9* [15, 21]	−23.3*	+5.6	+36.7*	+31.1*	+21.1*	+2.2	+52.2*	+26.7*	+8.9	+0.0	+35.6*
Combined	+23.6*	+17.7* [14, 21]	+14.4*	+5.6	+30.0*	+22.2*	+14.4	+14.4*	+10.0	+28.9*	+12.2*	+12.2*	+30.0*
Source citations	+21.1*	+18.2* [14, 23]	+22.2*	+5.6	+30.0*	+17.8*	+11.1	+20.0*	+16.7	+25.6*	+6.7	+13.3*	+31.1*
Statistics	+13.7*	+11.8* [8, 15]	+17.8*	+2.2	+13.3	+8.9	+11.1	+16.7*	+2.2	+20.0*	+6.7	+13.3*	+17.8*
Keyword stuffing	+12.4*	+8.0* [5, 11]	+3.3	+5.6	+8.9	+6.7	+14.4	+4.4	+11.1	+10.0	+2.2	+7.8	+13.3
Persuasive tone	+10.2*	+8.0* [5, 11]	−11.1	+4.4	+25.6*	+2.2	+2.2	+7.8	+11.1	+17.8*	+14.4*	−2.2	+15.6
Expert quotation	+8.1*	+3.6 [1, 6]	+1.1	+0.0	+8.9	−3.3	+5.6	+7.8	+0.0	+5.6	+0.0	+2.2	+12.2
Freshness date	+4.0*	−1.4 [−4, 1]	−3.3	+3.3	+1.1	−6.7	−7.8	+1.1	−1.1	+1.1	−1.1	−2.2	+0.0
Stated limitations	+3.9*	+3.2* [1, 5]	+0.0	+3.3	+11.1	−2.2	+2.2	+3.3	−2.2	+5.6	+7.8	+4.4	+2.2
Q&A format	+2.7	+2.6 [1, 5]	−3.3	+4.4	+8.9	+4.4	+2.2	+3.3	−3.3	+2.2	+8.9	+1.1	+0.0
Answer-first summary	+2.1	+1.5 [−1, 4]	−1.1	+0.0	+5.6	−1.1	+3.3	+3.3	−3.3	+8.9	+2.2	−1.1	+0.0
Spec table	+2.1	+0.6 [−2, 3]	−4.4	+0.0	+5.6	−11.1	+8.9	+4.4	−3.3	+3.3	−4.4	+0.0	+7.8
Boilerplate placebo	+1.8	+1.7 [0, 4]	−1.1	+1.1	+1.1	−4.4	+2.2	+4.4	+1.1	+3.3	+6.7	+2.2	+2.2
Baseline mention rate (%)	51	61	62	94	50	59	63	52	42	67	86	49	47

Sampling. At temperature 1, four models gave treatment effects of the same sign as at temperature 0 in 21 of 24 cases ($r = 0.96$), and no effect change survived the specified correction. This check does not establish equivalence; the largest difference was for citations (−6.9 points, [−12.8, −0.3]).

Absorption detector. For statistics and citations, the keyword detector agreed closely with both model judges (Cohen's $\kappa \geq 0.73$), and judges confirmed at least 94% of its positives (Table 14). It missed paraphrases, most often of stated limitations, while both judges also produced false positives. Their quotation labels are especially uncertain: they also marked baseline answers that never saw a quotation (35.0% and 40.0%). Detector-positive rates are descriptive, not one-sided bounds on accurate transmission. The judges marked statistics (69.2% and 67.5%) and citations (81.7% and 79.2%) more often than persuasive claims (12.5% and 21.7%).

Exploratory usage weighting. In a supplementary analysis only, developer API models were weighted by external estimates of assistant web-visit shares (Table 17). Third-party endorsement raised mentions by 31.6 points (95% CI [25.5, 37.9]), citations by 16.6 ([11.6, 21.7]), statistics by 10.7 ([6.3, 15.8]), keyword stuffing by 9.8 ([5.7, 14.1]), and persuasive copy by 7.2 ([3.3, 11.0]). These estimates depend on imperfect developer-model proxies; they are not estimates of effects in the named deployed assistants and are not principal evidence.

Table 13: Long pages and temperature 1: pooled effects on mentions (percentage points; 95% CI for the robustness version), and the change in each effect (robustness minus main, paired within trial). *Holm-adjusted $p < 0.05$ across the treatments in each check. "Main" uses the same models and trials as the check.

	Long pages			Temperature 1		
Treatment	Main	Robustness [95% CI]	Change	Main	Robustness [95% CI]	Change
Statistics	+13.8*	+6.4* [4, 9]	−7.5*	+13.3*	+11.4* [6, 17]	−1.9
Source citations	+21.2*	+11.8* [8, 16]	−9.4*	+25.0*	+18.1* [11, 25]	−6.9
Expert quotation	+7.1*	+2.3 [0, 5]	−4.8*	–	–	–
Third-party endorsement	+36.9*	+24.6* [21, 29]	−12.3*	+39.2*	+36.4* [30, 43]	−2.8
Persuasive tone	+10.2*	+0.1 [−2, 2]	−10.1*	–	–	–
Keyword stuffing	+11.7*	+21.8* [18, 26]	+10.1*	+15.3*	+10.0* [4, 16]	−5.3
Boilerplate placebo	+1.5	−0.4 [−2, 1]	−2.0	−0.3	+1.4 [−4, 7]	+1.7
Source-embedded instruction	+27.3*	+13.6* [11, 16]	−13.7*	+31.7*	+27.2* [21, 34]	−4.4

Table 14: The absorption detector against two model judges, for 720 answers that named the target (120 per edit) and 240 paired baseline answers that never saw the added content. "Repeated": share of answers labelled as passing on the added content (%); "False pos.": share of baseline answers labelled so (%); κ: Cohen's kappa between detector and judge. The judges agreed with each other at $\kappa = 0.84$.

	Keyword detector		DeepSeek V3.2			Claude Haiku 4.5		
Edit	Repeated	False pos.	Repeated	False pos.	κ	Repeated	False pos.	κ
Statistics	58	0	68	20	0.73	69	0	0.77
Source citations	76	0	79	15	0.81	82	0	0.83
Expert quotation	23	0	67	40	0.26	73	35	0.20
Stated limitations	12	0	32	0	0.37	43	5	0.31
Freshness date	1	0	7	2	0.21	2	0	0.66
Persuasive tone	19	0	22	0	0.72	12	0	0.44
All six	32	0	46	13	0.65	47	7	0.61

Live interface audit methods. The audit was specified in the experiment's documentation before the first query. The four local scenarios name invented towns that a live engine cannot

answer, so 26 questions were asked. Perplexity ran with default settings in its incognito mode for 22 questions; after a search limit forced a new sign-in, three were asked signed in without it and one while it was being switched back on (each record is flagged). Gemini ran in a temporary chat. Answers were read at human pace, one question at a time, and for each we kept the cited URLs (Perplexity) or source names (Gemini) and the brands in bold or in tables. A source counts as a vendor's when its owner makes or sells what that question asks about, including retailers and the vendor's own blog or documentation; the hand labels for every question are listed in the released `analyze_audit.py`. Robots.txt rules for each cited page's path were evaluated with Study 1's parser. ChatGPT was added after the Perplexity and Gemini results were known, and its procedure was written down before its first query: the same 26 questions in the same order, each in a fresh temporary chat set to "Unpersonalized" (which ignores memory and custom instructions), on the free plan, which offers no model choice, with web search left to ChatGPT. After a one-question pilot we recorded links that ChatGPT attaches to recommended products ("Official website") separately from cited sources; the pilot is kept. One answer that showed web images is recorded as such, without sources.

Live interface audit results. Perplexity displayed 270 citations on 172 domains; the hand-labelled third-party share was 73%, and vendors supplied most citations for only two questions. Gemini displayed 64 sources, of which 32 were vendor sources; this snapshot does not establish that third parties held a majority. Of Perplexity's displayed citations, 56 (21%, on 27 domains and in 19 questions) pointed to pages whose robots.txt, fetched the same day, disallowed `PerplexityBot`; 37 also disallowed `Perplexity-User`. A robots.txt rule governs crawling, not citation, and the observation cannot establish that the engine ignored the rule. ChatGPT gave product recommendations in 25 answers and attached official-site links in twelve answers (58 links); in three it asked about country or budget before a final pick, and once the interface showed web images. These observations do not reveal backend retrieval or establish that answers came solely from parametric memory. Several answers were tailored to India, and one referred to the account's approximate location.

J Searchable-corpus task

Design. Each scenario's pool holds the main design's five brand pages and buyer's guide and seven more fictional brands generated in the same way (names checked in DNS; one unedited paragraph each), 13 pages per scenario and 390 in one index. In each episode the target's page, or the guide in the third-party condition, is the version for that condition, and every other page is unedited. The search tool ranks all 390 pages by reciprocal rank fusion ($k = 60$) of Okapi BM25 ($k_1 = 1.2$, $b = 0.75$, over title and text, lower-cased, stop words removed, plural endings stripped) and the cosine similarity of 256-dimensional OpenAI text-embedding-3-small vectors, and returns the top five with title, URL and the page sentence that shares most words with the query (extended by the next sentence when shorter than 20 words, at most 40 words). The page tool returns a page's title, URL and full text. Models call the tools natively: the Azure deployments through Chat Completions and GPT-6.1 Sol through OpenAI's Responses API, because it accepts function tools with reasoning only there. After eight calls the tools are withheld and a message asks for the answer. Settings are those of the main runs: temperature 0 where accepted, low reasoning effort for the GPT models, and a 4,000-token limit per call. The system prompt was:

> You are an AI search assistant with two tools: search, which searches the web and returns the top five results with short snippets, and open_page, which returns the text of a web page.

> Search before you answer, and answer the user's question using only the pages you found, not your prior knowledge of products. Cite the pages you rely on by putting their URL in square brackets, such as [https://example.com/], at the end of the relevant sentence. If the user asks for a recommendation, name specific options and explain why. Keep the answer under 250 words.

The user turn is the scenario's question alone. Trials ran in a fixed random order, all seven conditions of a trial together, and each model had a spending cap; a model would have been reported only descriptively with fewer than 60 complete trials. The analyses are those of the main design: paired differences from baseline within a trial, cluster bootstrap intervals over scenarios, sign-flip tests and Holm correction over the six treatments within each model and outcome, and pooled estimates that weight models equally.

Process. Table 15 describes how the models searched. "No search" counts answers given without a single search; "unseen URL cited" counts answers citing a URL that no search returned and the model did not open; "no pool brand" counts answers naming none of the scenario's 12 brands, for example because they asked the user a question first.

Table 15: How the models searched in the searchable-corpus task: searches and pages opened per episode, and shares of episodes (%) answered without a search, citing a URL no search returned, reaching the call limit, or naming none of the pool's brands (all seven conditions).

Model	Trials	Searches	Pages opened	No search	Unseen URL cited	Call limit	No pool brand	Cost ($)
GPT-6.1 Sol	90	2.0	3.0	0	0	1	0	6.15
GPT-5 mini	90	2.7	1.6	0	0	5	14	1.76
MAI-Thinking-1	90	0.7	0.6	29	29	0	29	8.16
DeepSeek V4 Pro	90	3.1	3.4	0	0	35	4	10.75

Table 16: Searchable-corpus task: effects on mentions with the pages given (main design, same models and trials) and when the model searches, per model, and effects on whether the target page appeared in search results and was opened (percentage points; * Holm-adjusted $p < 0.05$).

	Mentions						Target page	
Treatment	Pages given	Model searches [95% CI]	*GPT-6.1 Sol*	*GPT-5 mini*	*MAI-Thinking-1*	*DeepSeek V4 Pro*	In results	Opened
Keyword stuffing	+10.3	+17.2* [9, 26]	+14.4	+17.8*	+3.3	+33.3*	+54.4*	+26.1*
Third-party endorsement	+35.3	+13.3* [8, 19]	+4.4	+21.1*	+0.0	+27.8*	+2.8	+6.4
Answer-first summary	+3.1	+7.2* [2, 12]	+3.3	+11.1	+5.6	+8.9	+26.4*	+5.3
Source-embedded instruction	+23.9	+0.6 [−5, 6]	−2.2	+3.3	−1.1	+2.2	+7.2	+1.1
Source citations	+23.1	−1.1 [−5, 3]	+1.1	+4.4	−7.8	−2.2	+3.9	−1.1
Statistics	+14.7	−5.8 [−11, −1]	−4.4	−6.7	−7.8	−4.4	−2.2	−1.7
Baseline rate (%)	53	22	20	22	19	28	38	14

Exploratory usage-weighting proxies. The table reports the external web-visit shares used only for the supplementary exploratory weighting analysis. Shares are not measures of answer-

engine use, and the listed developer API models are imperfect proxies for deployed assistants; the table does not establish effects for the most-used production assistants.

Table 17: Exploratory usage-weighting proxies only. Shares are estimates of worldwide generative-AI web visits in May 2026 [38], not answer-engine use; the remaining share is not broken down by assistant. The listed developer API models are imperfect proxies, not the assistants' deployed models or orchestration.

Assistant (developer)	Web visits	Main analysis	Replication
ChatGPT (OpenAI)	≈53%	gpt-oss-120b°	GPT-5 mini, GPT-6.1 Sol
Gemini (Google)	27–28%	Gemma 3 27B°	Gemini 3.5 Flash-Lite
Claude (Anthropic)	≈9%	Claude Sonnet 4.6	Claude Opus 4.6, Claude Sonnet 5.5
Copilot (Microsoft)	together ≈10%	–	MAI-Thinking-1
Grok (xAI)		–	Grok 4.6
DeepSeek		DeepSeek V3.2	DeepSeek V4 Pro
Meta AI (Meta)		Llama 4 Maverick	–
Kimi (Moonshot AI)		–	Kimi K3, Kimi K3.1 Pro
Qwen (Alibaba)		Qwen3 32B, Qwen3 Next 80B	–
Le Chat (Mistral AI)		Mistral Large 3	–
Other developers		Nova Pro, Nova 2 Lite, MiniMax M2.5, Nemotron 3 Super	GLM 5

K Evidence for common AEO techniques

To set our results against prior evidence we grade AEO techniques from a structured narrative review (not a systematic review) of 2023–2026 papers, operator documentation and industry studies that disclose their methods. The grade measures only the strength of evidence that a technique matters for access or changes the measured outcome, not whether it is advisable; a separate column records the risks and the limits of our evidence. Grade A: a necessary condition documented by operators or standards, or consistent effects in at least two controlled studies. B: one controlled study, several that agree in direction but not in magnitude, or a large observational study. C: mixed, mechanism-only or uncontrolled evidence. D: controlled evidence of no effect. Table 18 lists the grades after adding this paper's results, which come from fictional brands, synthetic pages and API models, not deployed engines.

Table 18: Evidence for common AEO techniques, combining prior work with this paper. The grade (A–D, defined above) rates the strength of evidence only; risks and scope limits are listed separately. Effects are on mentions with the pages supplied, pooled over the 12 hosted models, unless stated; Sections 5.2 and 5.3 give the replication, long-page and search-task results.

Technique (stage)	Prior evidence	This paper	Grade	Risks and limits
Allow each engine's search crawler in robots.txt (access)	Operators document separate tokens [2, 32, 34]	9.3% of organisational sites disallow one by name	A	Governs one crawler's route; 21% of Perplexity's citations went to disallowed pages
Keep robots.txt reachable (access)	A 5xx response means full disallow [20]	3.3% unreachable	A	One snapshot; errors may be transient
Admit verified AI crawlers at the edge (access)	Default blocking by a major CDN [6]	18.6% of research-client requests refused or challenged	A	Our client was not a verified AI crawler
Serve content without JavaScript (access)	Most AI crawlers do not render [40]	11.2% of rendered homepages met the word-count proxy	B	Proxy only; text loss not shown to matter
Publish llms.txt (access)	Files are rarely read [24]	23.0% of sites publish one	C	Use by engines not tested
Third-party endorsement (selection)	AI search favours earned media [5]	+37.4 points, all models; holds in replication and search; smaller on long pages	B	One synthetic guide; adds information as well as a source
Statistics (selection)	Gains in GEO [1], diluted under competition [36]	+13.7 points, beats baseline and placebo; repeated in answers; none in search	B	Synthetic, unverified claims; a package, not a format
Source citations (selection)	As for statistics [1, 36]	+21.1 points, strongest single edit; repeated in answers; none in search	B	As for statistics
Expert quotation (selection)	Largest gain in GEO [1]	+8.1 points, beats baseline and placebo	B	The quotation includes a recommendation
Answer-first summary, Q&A, tables (selection)	Gains in citation [45]; formatting did little [41]	+2.1, +2.7, +2.1 points, none significant; the summary helped pages be found in search	C	One wording each
Freshness date (selection)	Recent dates helped [41]	+4.0 points, not above placebo; almost never repeated	C	One wording
Stated limitations (selection)	None	+3.9 points, not above placebo	C	One wording
Persuasive copy (selection)	Mixed [1, 36]	+10.2 points; rarely repeated; none on long pages	C	Superlatives can mislead users
Keyword stuffing (selection, retrieval)	No gain [1]	+12.4 points, beats placebo; +21.8 on long pages; raised retrieval in our retriever	C	Violates search spam policies [13]; our retriever has no spam filter
Instructions embedded in pages (selection)	Steer LLM recommendations and a commercial search product [21, 39]	+28.3 points, 10 of 12 models, rarely disclosed; none in search	B	Deceptive and against spam policies; never use

L Sensitivity analyses

The analyses in this appendix were added after the results were known and are exploratory (Appendix E).

Tranco rank bands. Table 19 shows, for each rank band, how many domains were drawn, retained and excluded, and the share of retained domains whose robots.txt disallows at least one answer agent. The shares describe the retained domains and are not weighted to the population of ranked domains. The number of exclusions in each band is known; the reason for each exclusion

was not recorded.

Table 19: Tranco sample by rank band: domains drawn, retained and excluded, and the share of retained domains whose robots.txt disallows an answer agent (95% Wilson intervals).

Tranco rank band	Drawn	Retained	Excluded	Disallows an answer agent (%) [95% CI]
1–1,000	250	148	102	10.8 [6.8, 16.8]
1,001–10,000	250	187	63	8.0 [4.9, 12.8]
10,001–100,000	250	178	72	2.8 [1.2, 6.4]
100,001–1,000,000	250	190	60	8.4 [5.2, 13.2]

Developer weighting and absolute rates. The main models come from 10 developers and the replication models from 8. The main analysis gives each model equal weight; Table 20 also gives each developer equal weight, averaging its models first, and recomputes that estimate leaving out one developer at a time. Neither weighting represents a population of models. Endorsement had the largest effect under both weightings and whichever developer was left out. Absolute rates show how much room each model had to improve: Claude Sonnet 5.5 named the target in 94.4% of baseline answers, so it could gain at most 5.6 points; per-model rates are released with the data. The boilerplate placebo is not an exact length control, but it is close: it added a median of 32 words to the target's page, within the range of the single edits (medians of 8 to 44 words); only the combined package added far more (149).

Table 20: Absolute mention rates (%) and paired effects (percentage points) for third-party endorsement, source citations, statistics and keyword stuffing. Both weightings describe the selected API models, not a population of models. The last column is the range of the developer-weighted effect when one developer is left out, a sensitivity analysis rather than a confidence interval.

		Equal weight per model			Equal weight per developer		Leaving out
Models	Treatment	Baseline	Treated	Effect	Baseline/treated	Effect [95% CI]	one developer
Main	Third-party endorsement	51.0	88.4	+37.4	52.6/89.9	+37.3 [32.7, 42.0]	+35.8 to +38.3
Main	Source citations	51.0	72.1	+21.1	52.6/74.0	+21.4 [17.5, 25.6]	+20.7 to +22.2
Main	Statistics	51.0	64.7	+13.7	52.6/65.8	+13.3 [9.3, 17.1]	+12.8 to +14.1
Main	Keyword stuffing	51.0	63.4	+12.4	52.6/64.7	+12.2 [9.3, 15.0]	+10.6 to +13.5
Replication	Third-party endorsement	61.0	91.5	+30.5	58.5/91.4	+32.9 [28.1, 38.1]	+31.0 to +35.1
Replication	Source citations	61.0	79.2	+18.2	58.5/78.5	+20.1 [15.4, 24.9]	+18.5 to +21.5
Replication	Statistics	61.0	72.8	+11.8	58.5/70.5	+12.0 [8.3, 16.0]	+10.9 to +13.4
Replication	Keyword stuffing	61.0	69.0	+8.0	58.5/67.1	+8.6 [5.0, 12.6]	+7.9 to +9.2

Search-task retrieval and model sensitivity. Table 21 replays each model's baseline search queries against the pool as edited in each condition, without calling any model, and records whether the target's page ranked in the top five under BM25 alone, embeddings alone and their fusion. Queries were available from 70 trials for MAI-Thinking-1, which often answered without searching, and from all 90 trials for every other model; pooled estimates weight models equally. Citations changed the fused top-five rate by +2.86 points (Holm $p = 0.840$). The last column repeats the mention analysis without MAI-Thinking-1. These are sensitivity analyses of the tested retriever, not an experiment that isolates retrieval. Per-model rates of searching, snippet exposure, page opening, mention, first mention and citation are released with the data.

Table 21: Search task: effects on the share of replayed baseline queries that returned the target's page in the top five, by ranking method, and effects on mentions without MAI-Thinking-1 (percentage points; 95% scenario-clustered intervals; *Holm-adjusted $p < 0.05$ within each column).

	Target page in top five			Mentions without
Treatment	BM25	Embedding	Fused	MAI-Thinking-1 [95% CI]
Third-party endorsement	-1.0	0.0	-0.1	$+17.8$ $[11.9, 23.7]^*$
Source citations	$+0.1$	$+3.6$	$+2.9$	$+1.1$ $[-3.3, 5.6]$
Statistics	-2.9	-2.7	-4.4	-5.2 $[-10.7, 0.0]$
Keyword stuffing	$+64.5^*$	$+54.4^*$	$+64.0^*$	$+21.9$ $[13.0, 30.7]^*$
Answer-first summary	$+40.3^*$	$+1.7$	$+24.8^*$	$+7.8$ $[3.3, 12.6]^*$
Source-embedded instruction	$+2.1$	$+3.6^*$	$+5.1$	$+1.1$ $[-4.4, 6.7]$